



Software Design & Software Architecture



Software architecture, definition (1)

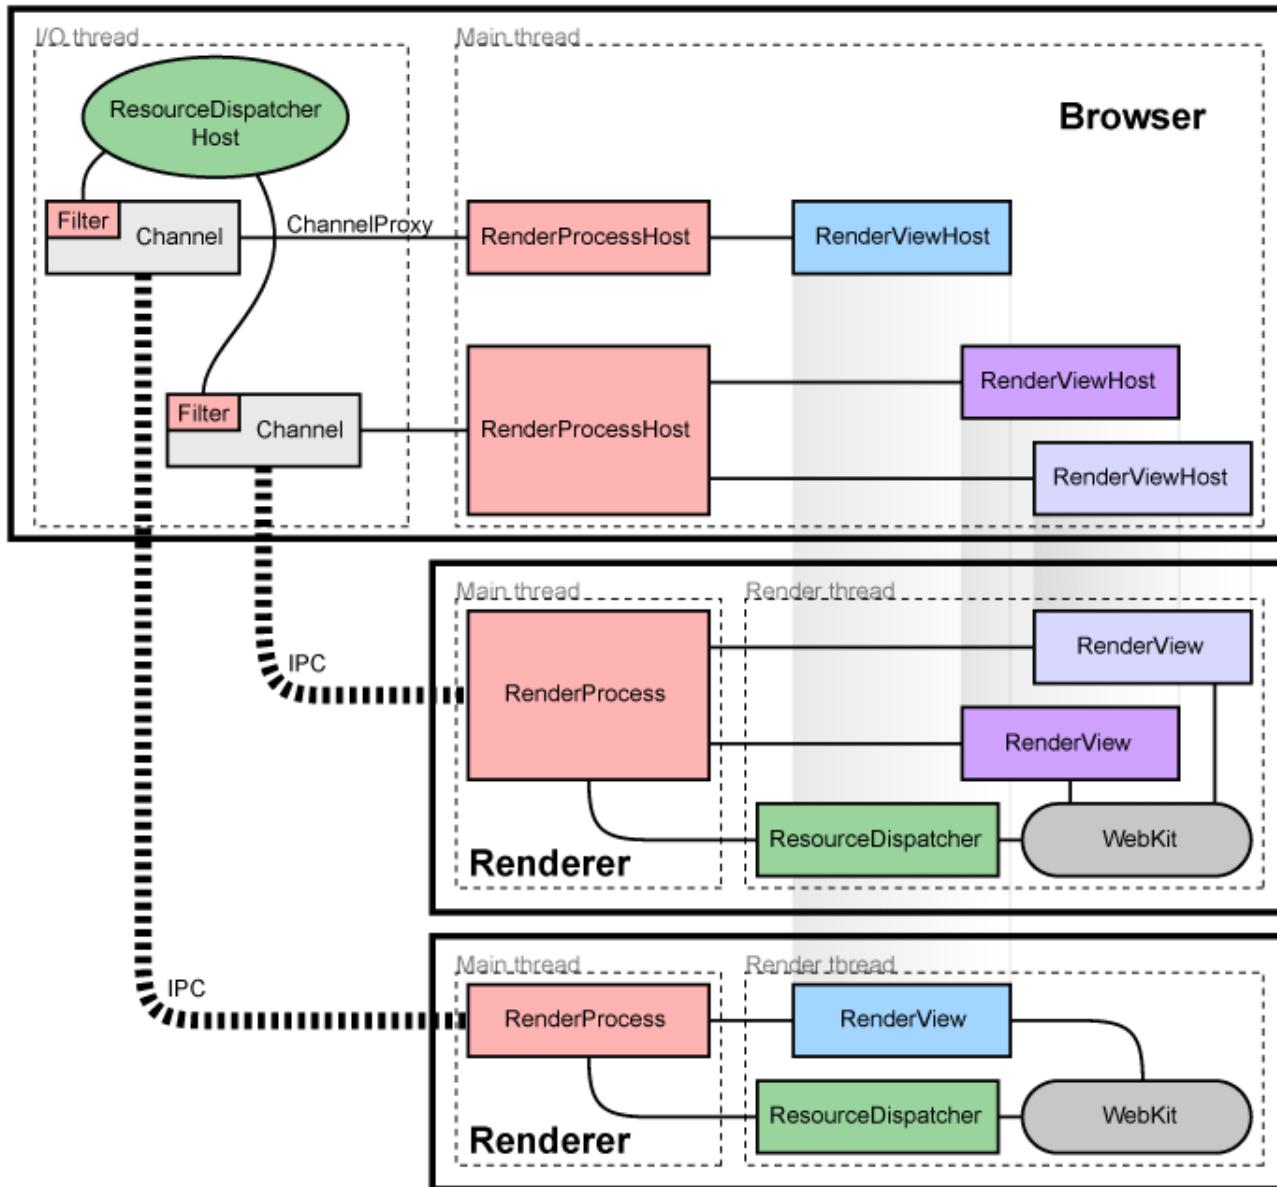
“The architecture of a software system defines that system in terms of computational components and interactions among those components.”

[Shaw and Garlan, '96], Shaw and Garlan, *Software Architecture, Perspectives on an Emerging Discipline*, Prentice-Hall, 1996.



Example: Chromium Architectural overview

<https://www.chromium.org/developers/design-documents/multi-process-architecture>



architecture

Software Architecture, definition (2)



“The software architecture of a system is the structure or structures of the system, which comprise software elements, the externally visible properties of those elements, and the relationships among them.”

[Clements and Kazman, 2003] Clements and Kazman, *Software Architecture in Practice*, SEI Series in Software Engineering. Addison-Wesley, 2003.



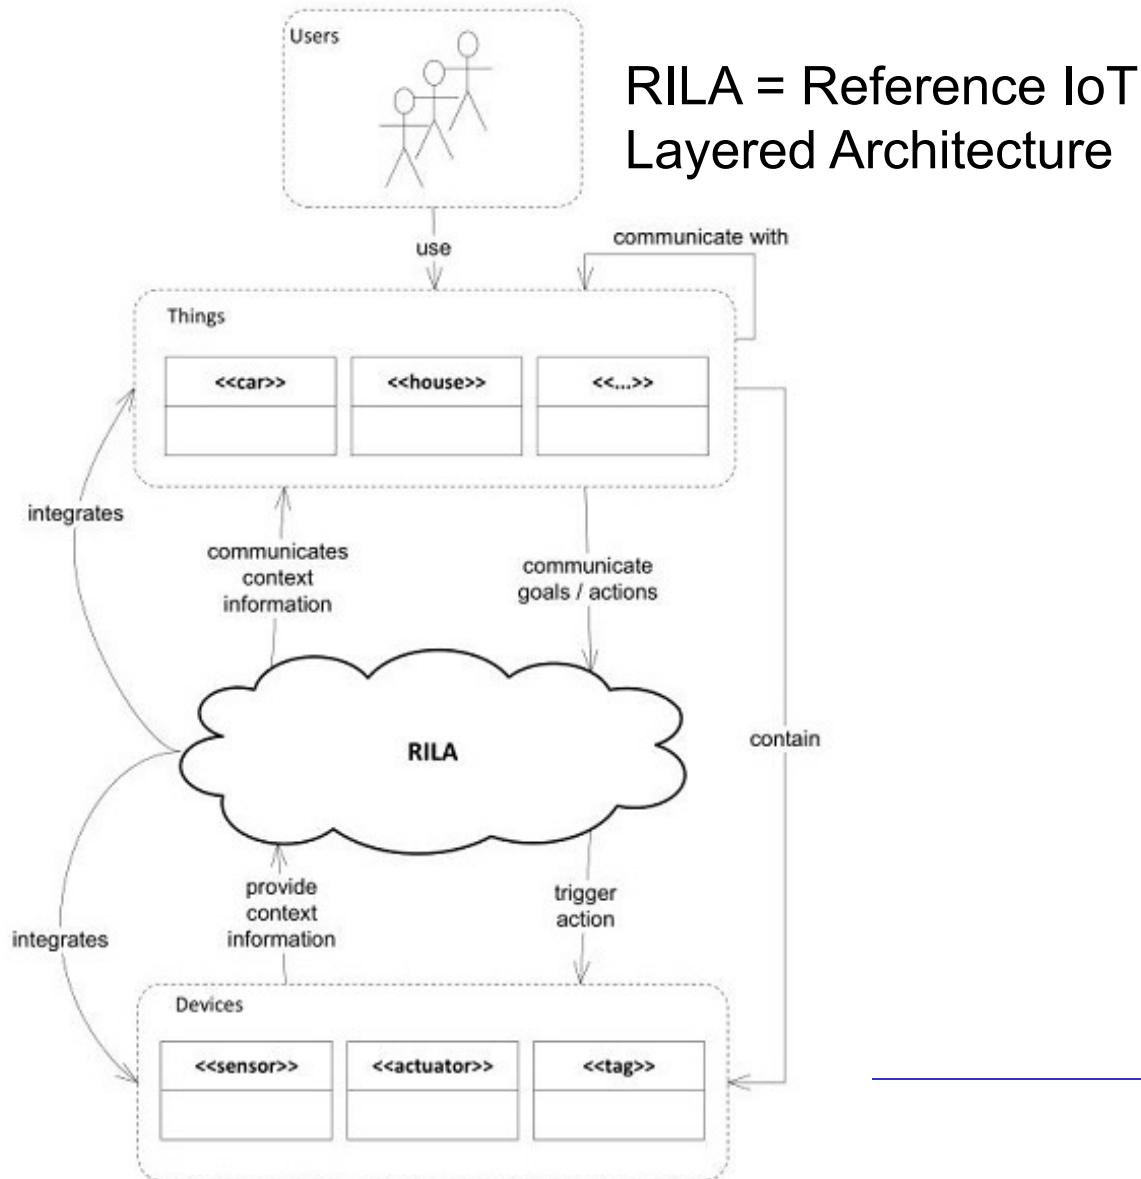
Software Architecture

- Important issues raised in this definition:
 - ▶ multiple system structures;
 - ▶ externally visible (observable) properties of components.
- The definition does not include:
 - ▶ the process;
 - ▶ rules and guidelines;
 - ▶ architectural styles.

Example: an Internet of Things reference architecture (1)

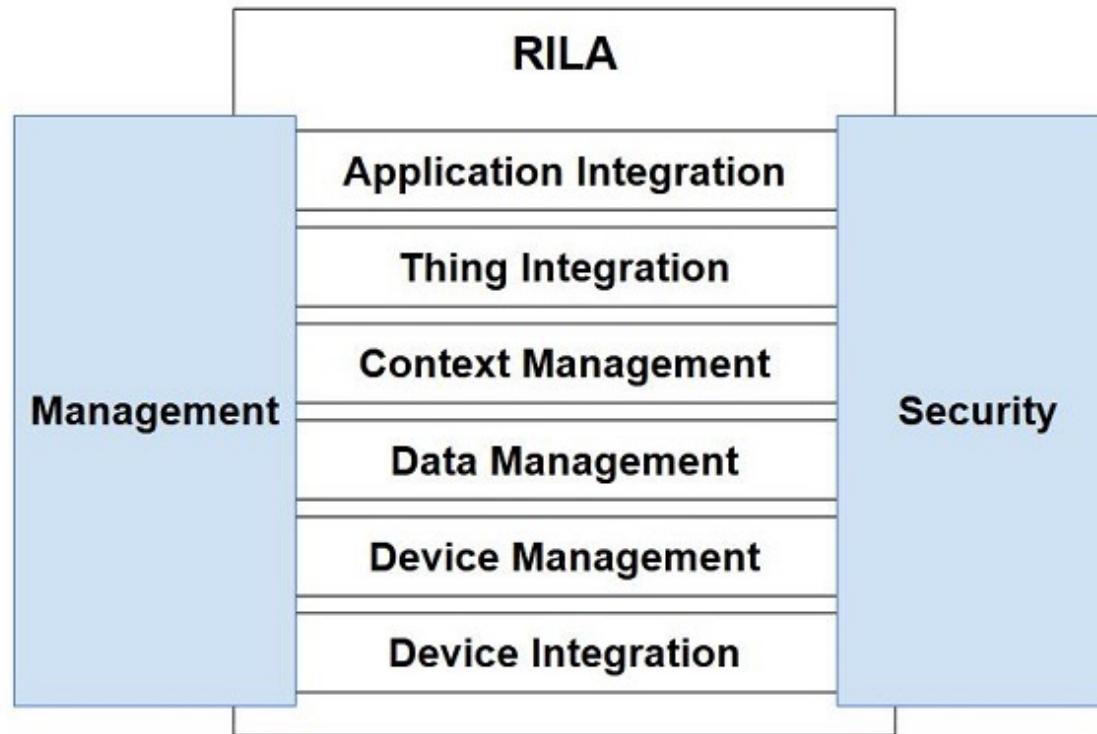


<https://www.infoq.com/articles/internet-of-things-reference-architecture>



Example: an Internet of Things reference architecture (2)

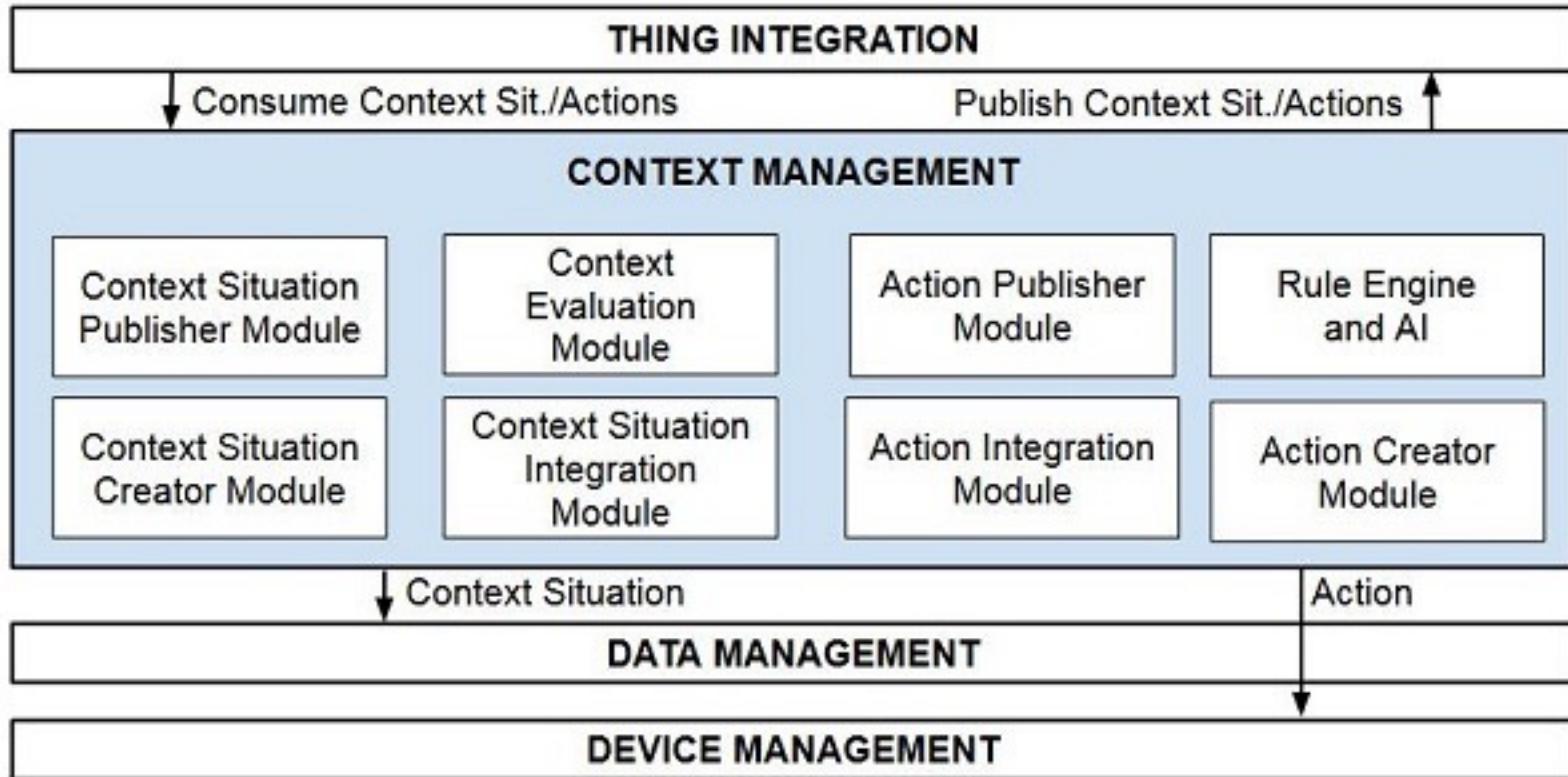
<https://www.infoq.com/articles/internet-of-things-reference-architecture>



Example: an Internet of Things reference architecture (3)



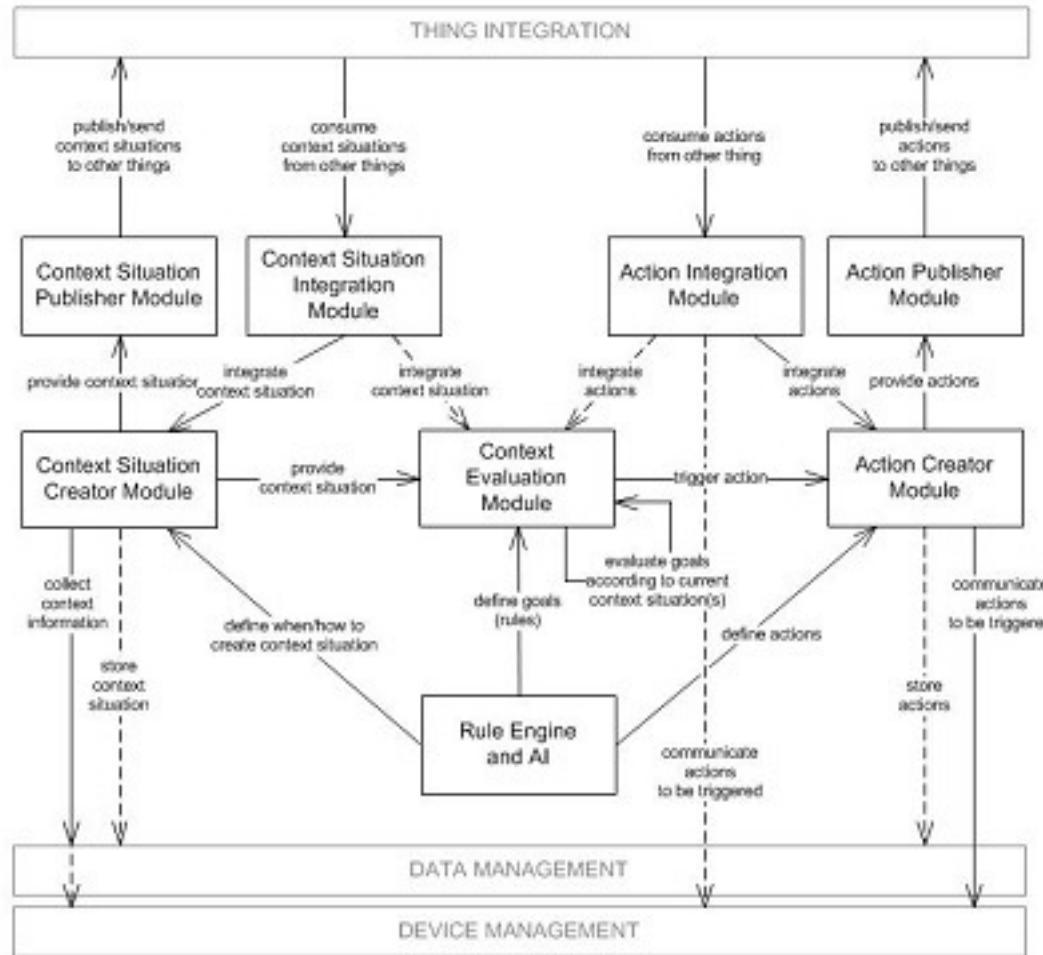
<https://www.infoq.com/articles/internet-of-things-reference-architecture>



Example: an Internet of Things reference architecture (4)



<https://www.infoq.com/articles/internet-of-things-reference-architecture>



What can we understand from this layered architecture?



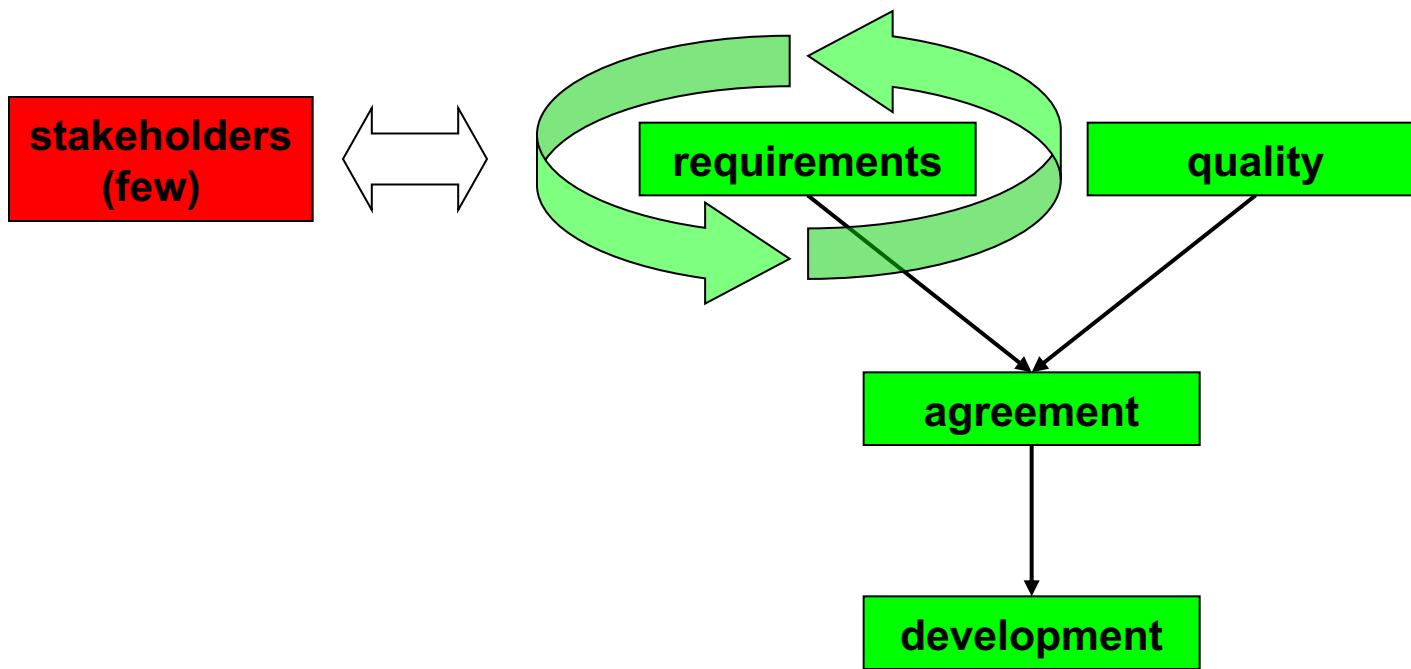
- Layers enable separation of concerns
- We can identify the main focuses of each layer at a glance.
- Each layer can be implemented using a different team
- Each layer remains independent from the others, provided that it respects the protocols defined for communication with the neighbour layers



Why Is Architecture Important?

- Architecture is the vehicle for communication: internal (different teams), external (teams and stakeholders)
- Architecture manifests the earliest set of design decisions
 - ▶ Introduces constraints on implementation
 - ▶ Introduces organizational structure
 - ▶ Inhibits or enables quality attributes
- Architecture is a transferable abstraction of a system
 - ▶ Product lines share a common architecture
 - ▶ Allows for template-based development
 - ▶ Basis for training

Pre-architecture life cycle

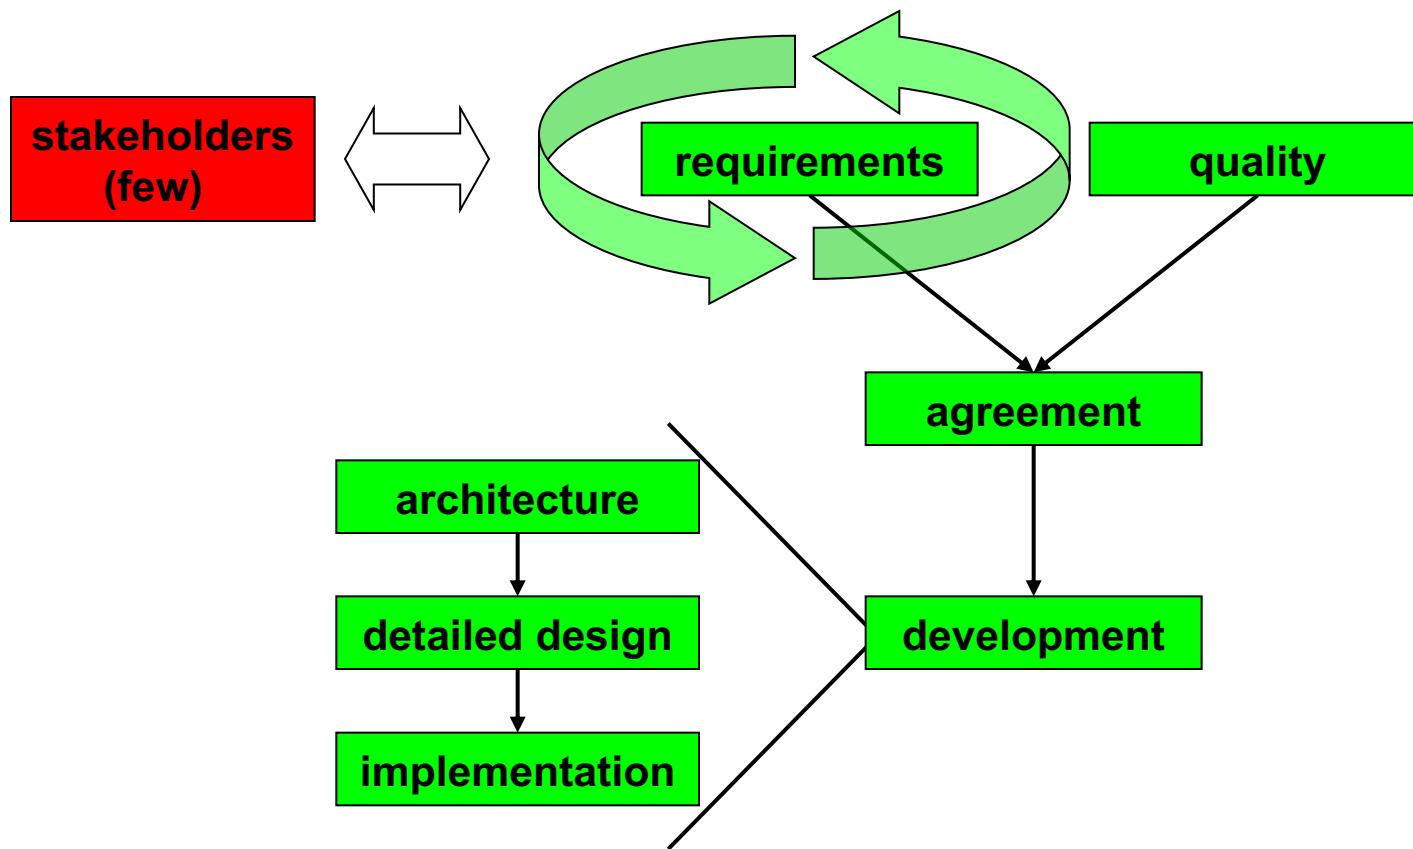




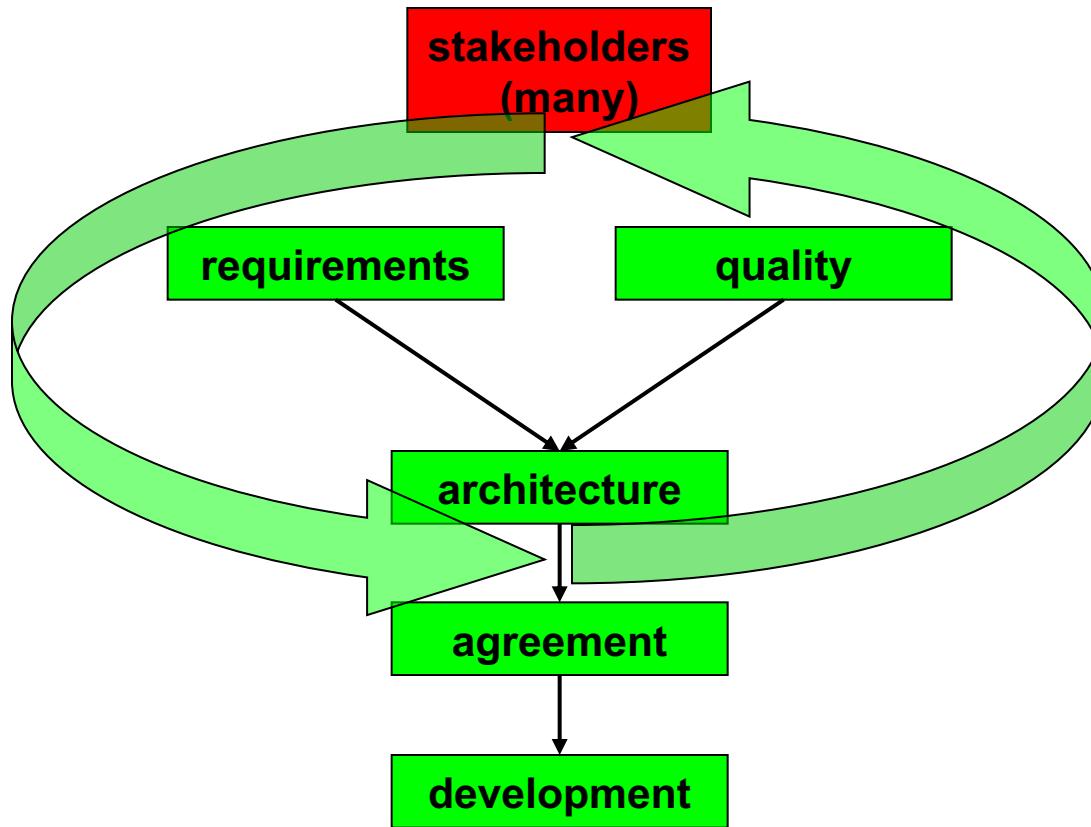
Characteristics

- Iteration mainly on functional requirements
- Few stakeholders involved
- No balancing of functional and quality requirements

Adding architecture, the easy way



Giving a more crucial role to the architecture





Characteristics

- Iteration on both functional and quality requirements
- Many stakeholders involved
- Balancing of functional, quality requirements, and architectural aspects



The architecture is intrinsic to our software!

A simple example: a trivial prototype of a “banking application” - Account



```
import java.rmi.*;
public interface Account extends Remote {
    public float balance() throws RemoteException;
}

import [...]
public class AccountImpl extends UnicastRemoteObject
    implements Account {
    // implementation of class (method balance, etc.)
}
```

A simple example: a trivial prototype of a “banking application” - AccountFactory



```
import java.rmi.*;  
public interface AccountFactory extends Remote {  
    public Account createAccount(String userName, int balance)  
        throws RemoteException;  
    public int getLoad() throws RemoteException;  
}  
  
import [...]  
public class AccountFactoryImpl extends UnicastRemoteObject  
    implements AccountFactory {  
    // implementation of class  
    // (methods createAccount, getLoad, etc.)  
    // ...  
}
```

A simple example: a trivial prototype of a “banking application” - AccountManager



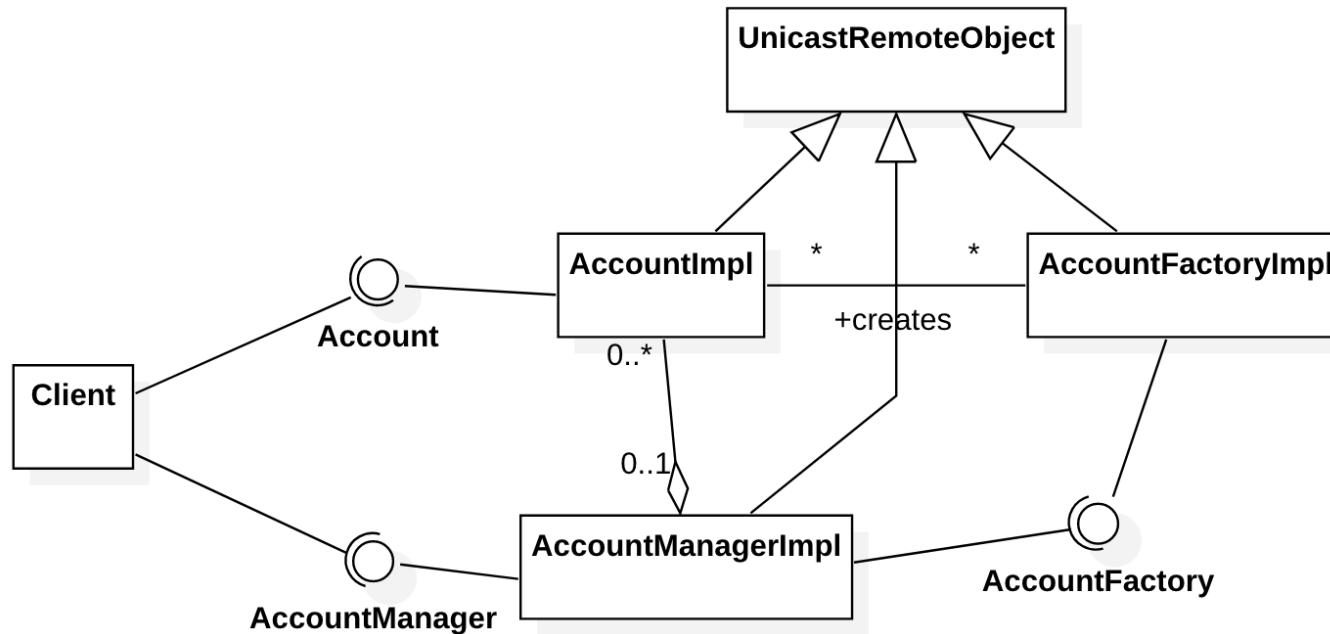
```
import java.rmi.*;  
public interface AccountManager extends Remote {  
    public Account open(String name) throws RemoteException;  
}  
  
import [...]  
public class AccountManagerImpl extends UnicastRemoteObject  
    implements AccountManager {  
  
    private Map<...> accounts = new HashMap<>();  
    // implementation of class (attributes, method open, etc.)  
    public void addFactory(String hostName) { ... }  
    public AccountFactory selectFactory() { ... }  
}
```

A simple example: a trivial prototype of a “banking application” - Client



```
import java.rmi.*;
public class Client {
    public static void main (String[] args) {
        try {
            // ... some stuff
            AccountManager manager =
                (AccountManager) Naming.lookup( "//" + args[0] +
                                                "/AccountManager");
            String name = "Jack B. Quick";
            Account account = manager.open(name);
            float balance = account.balance();
            System.out.println ("The balance in " + name +
                                "'s account is $" + balance);
        } catch (Exception e) { e.printStackTrace(); }
    }
}
```

A simple example: a trivial prototype of a “banking application” - Low level design view



- What info do I get from this?
 - ▶ Hierarchy
 - ▶ The client uses remote objects (sends messages to objects that extend **UnicastRemoteObject**)
 - ▶ Client uses two interfaces
 - ▶ The other interface is used by **AccountManager**

What happens if changes are applied on the source code?



- Architectural degradation
 - ▶ Changes in the software system are not reflected in the corresponding architectural description
 - ▶ Misalignment between documentation and implementation
 - Why does this happen?
 - ▶ Lack of a documented architecture
 - ▶ Inadequate processes and supporting tools
 - ▶ Developer sloppiness (or technical debt)
 - ▶ Perception of short deadlines
-

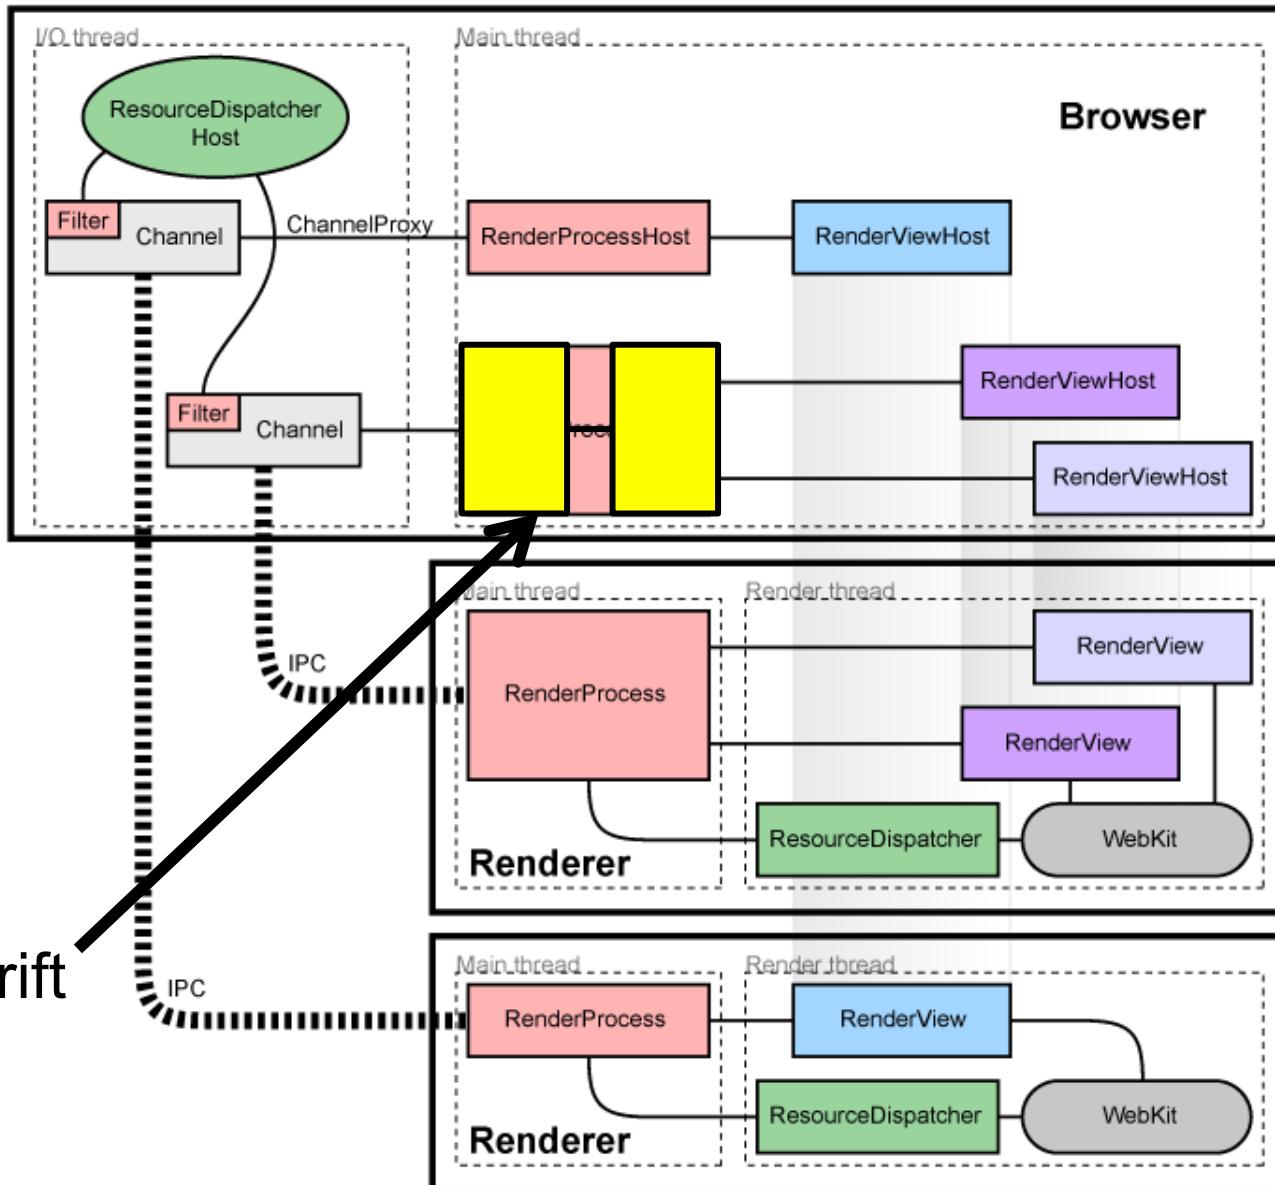


Drift and erosion

- Two kinds of degradation
 - ▶ **Drift:** the change is not included nor implied by the architecture but it does not necessarily imply violations
 - ▶ **Erosion:** the change violates prescriptive design decisions
-

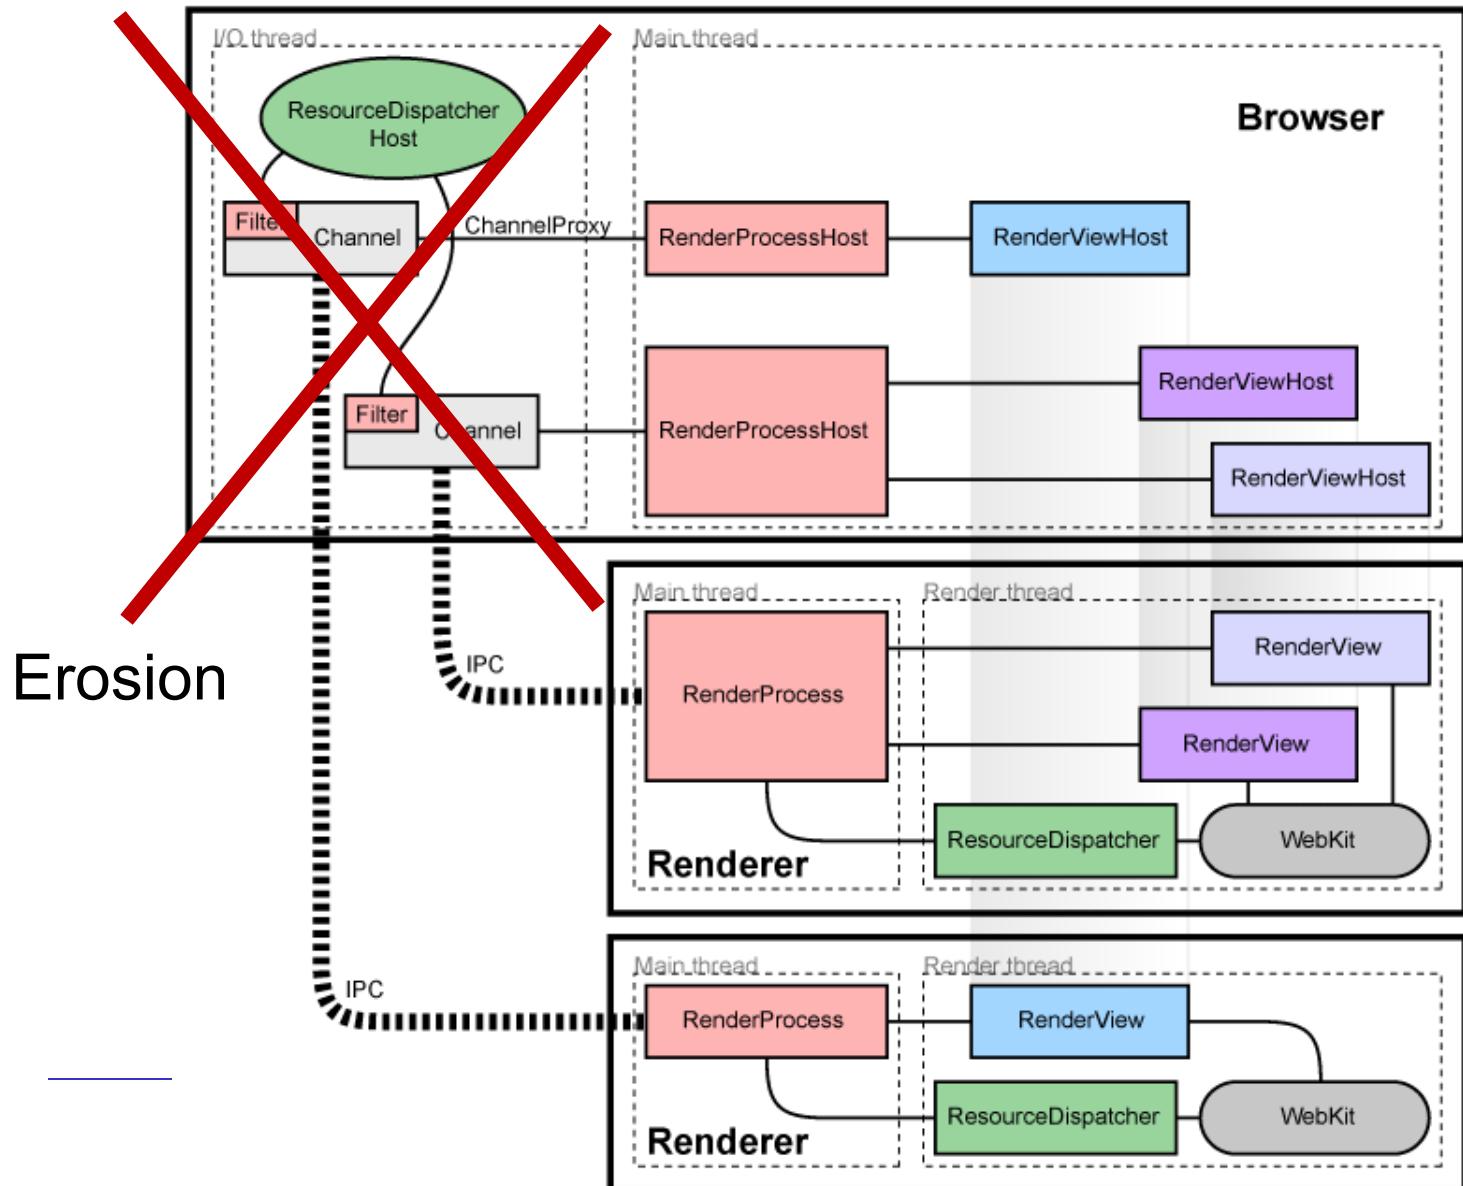
Example – Chromium Architecture

<https://www.chromium.org/developers/design-documents/multi-process-architecture>



Example – Chromium Architecture

<https://www.chromium.org/developers/design-documents/multi-process-architecture>





- The notion of quality is central: the SA is derived according to qualities of the system at the earliest possible stage.
- External qualities: performance, security, availability, functionality, usability
- Internal qualities: modifiability, portability, reusability, integrability, testability



Architectural styles

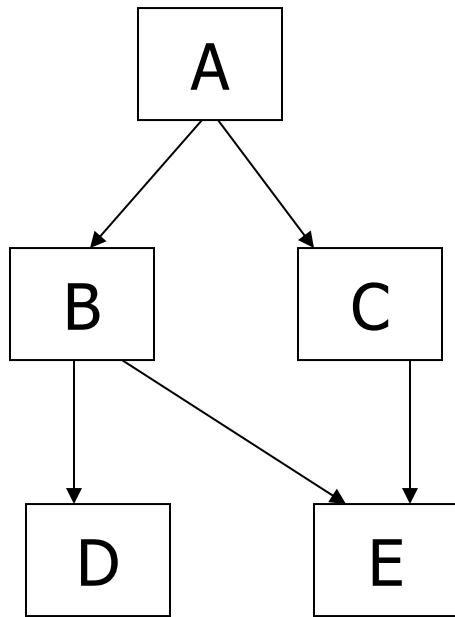


Architectural Style

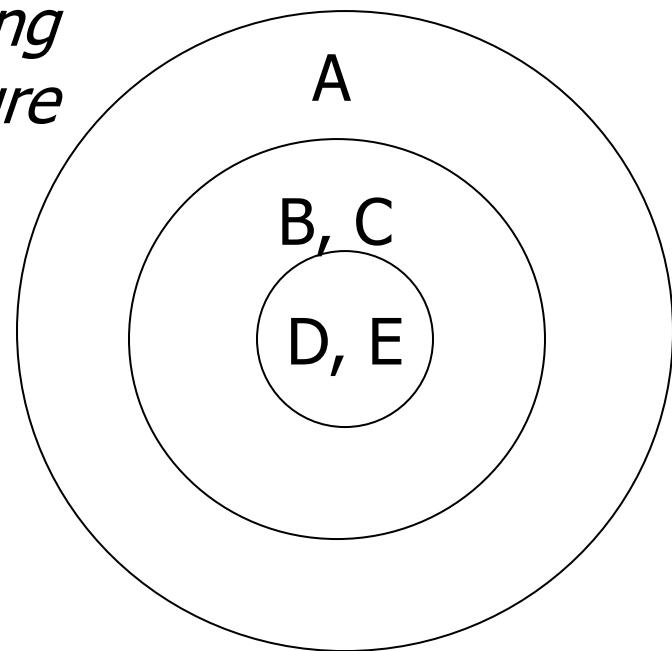
- “an architectural style determines the **vocabulary** of components and connectors that can be used in instances of that style, together with a **set of constraints** on how they can be combined. These can include **topological** constraints on architectural descriptions (e.g., no cycles). Other constraints—say, having to do with **execution** semantics—might also be part of the style definition.” [Garland & Shaw]
- [Garland & Shaw] David Garlan and Mary Shaw, “*An Introduction to Software Architecture*”, Jan 1994, CMU-CS-94-166
 - ▶ http://www.cs.cmu.edu/afs/cs/project/able/ftp/intro_softarch/intro_softarch.pdf

Layered style

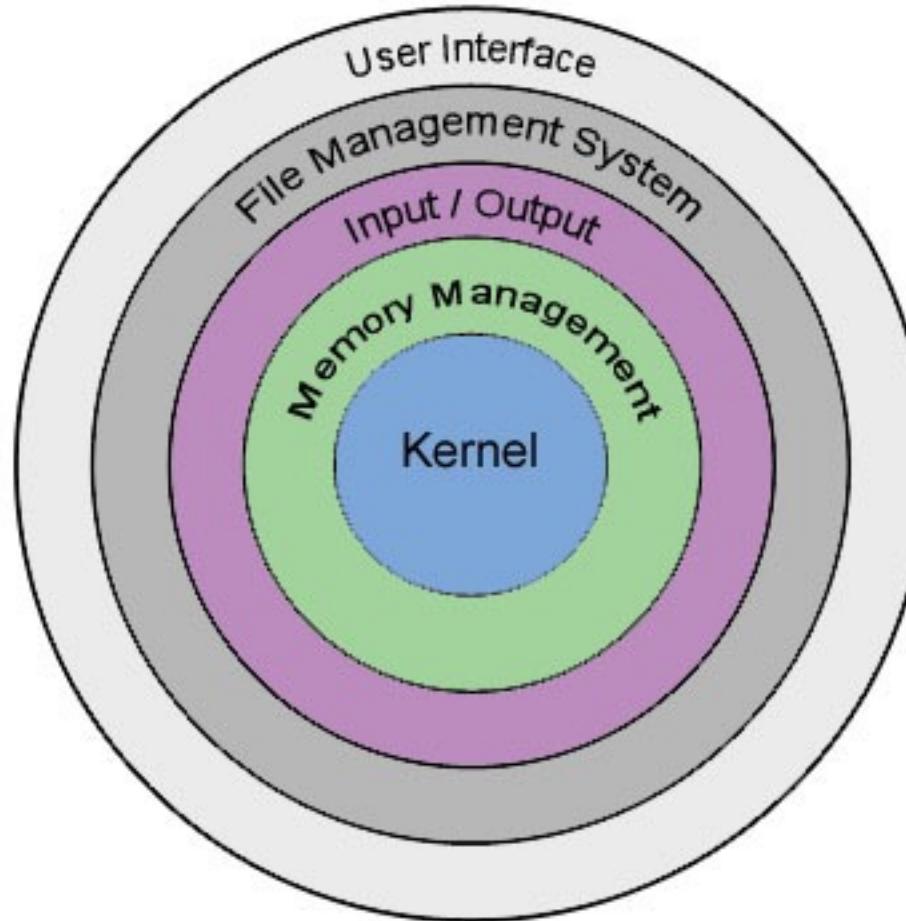
- The system is organized through **abstraction levels**, as a hierarchy of abstract machines
- **Hierarchy** is given by the **USE** relation



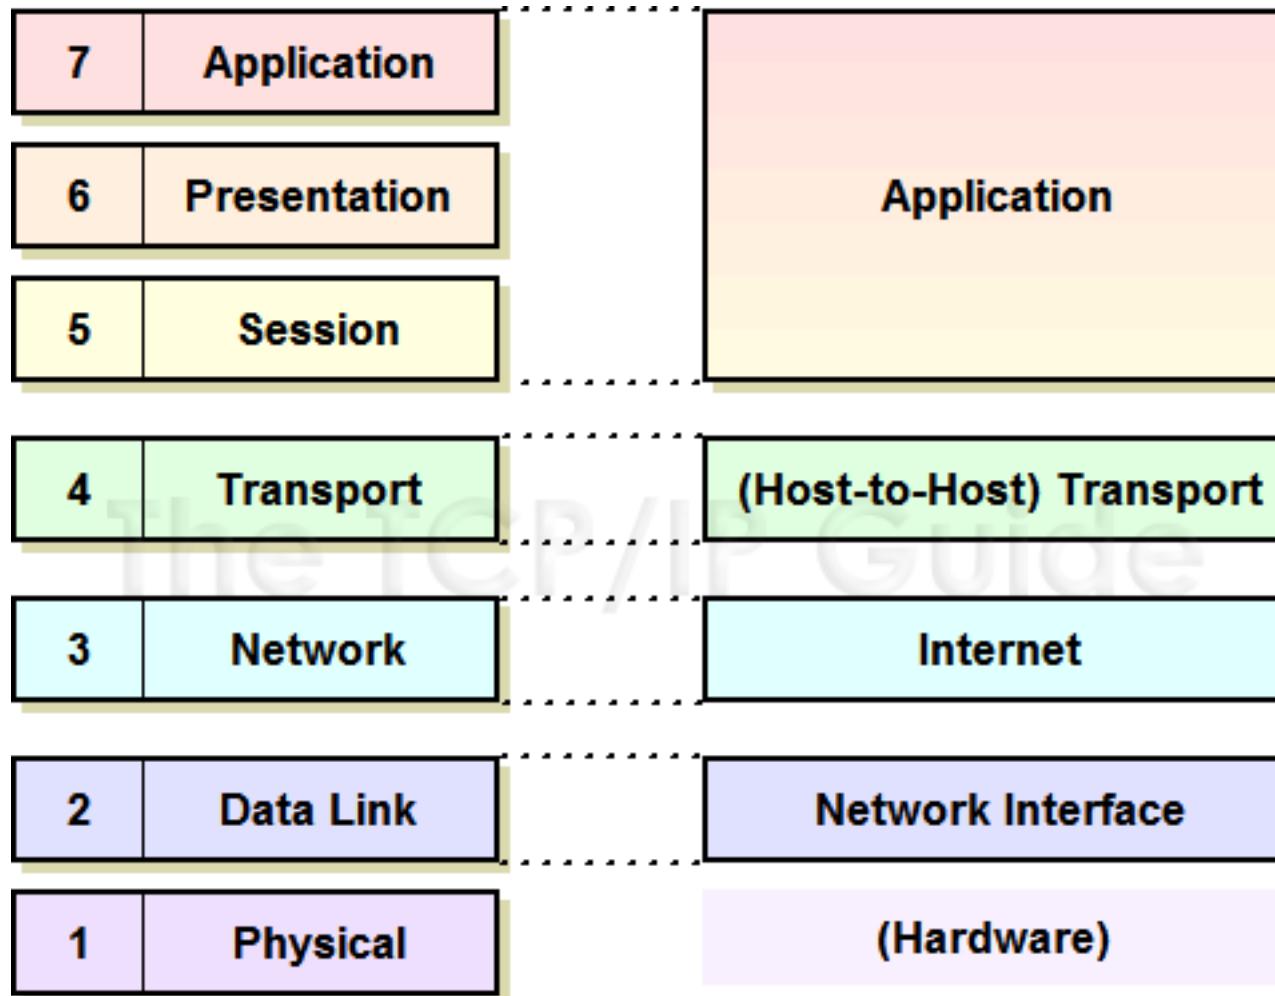
*onion-ring
structure*



Well-known layered systems: An Operating System



Well-known layered systems: The OSI and TCP/IP models



OSI Model

TCP/IP Model



When to use a layered architecture

- Abstraction layers to dominate complexity
- When it is possible to identify a specific (bounded) concern for each layer
- When it is possible to identify clearly the communication protocols between layers

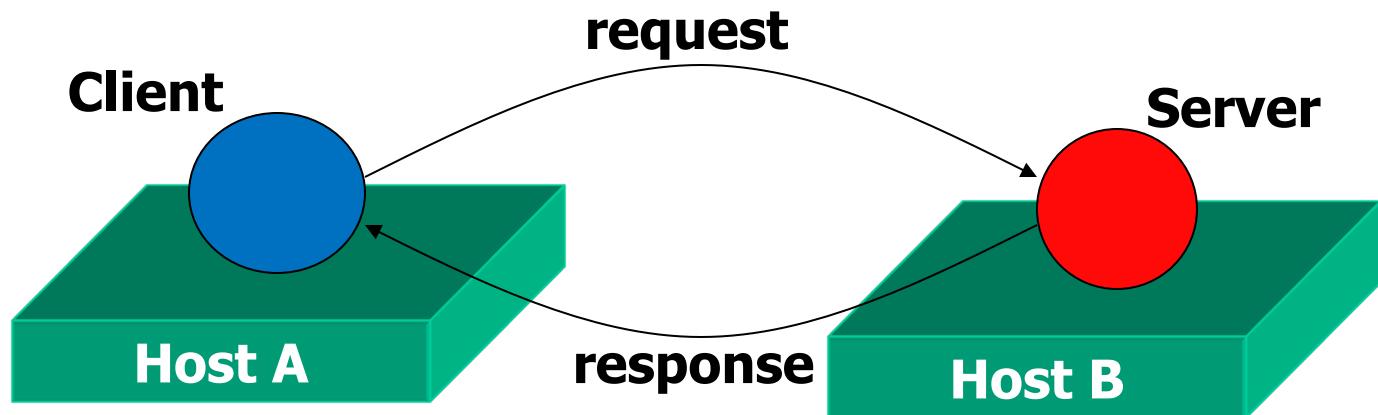


Client/Server

- Well-known architectural style for distributed applications
- Client and server are decoupled
 - ▶ Different processes with well-defined interface
 - ▶ Server is accessible only through interface
 - ▶ Server is unaware of specific client instances
 - ▶ Client and server have independent lifecycles
 - ▶ Client and server are defined by a set of hw/sw modules
 - ▶ May be distributed or run in the same compute unit
- Client and server play different roles
 - ▶ Server exports a set of services
 - ▶ Client uses the services and initiates communication
 - ▶ Communication through messages or through remote invocations
 - ▶ A client may be server for other components



Client /Server: a general schema

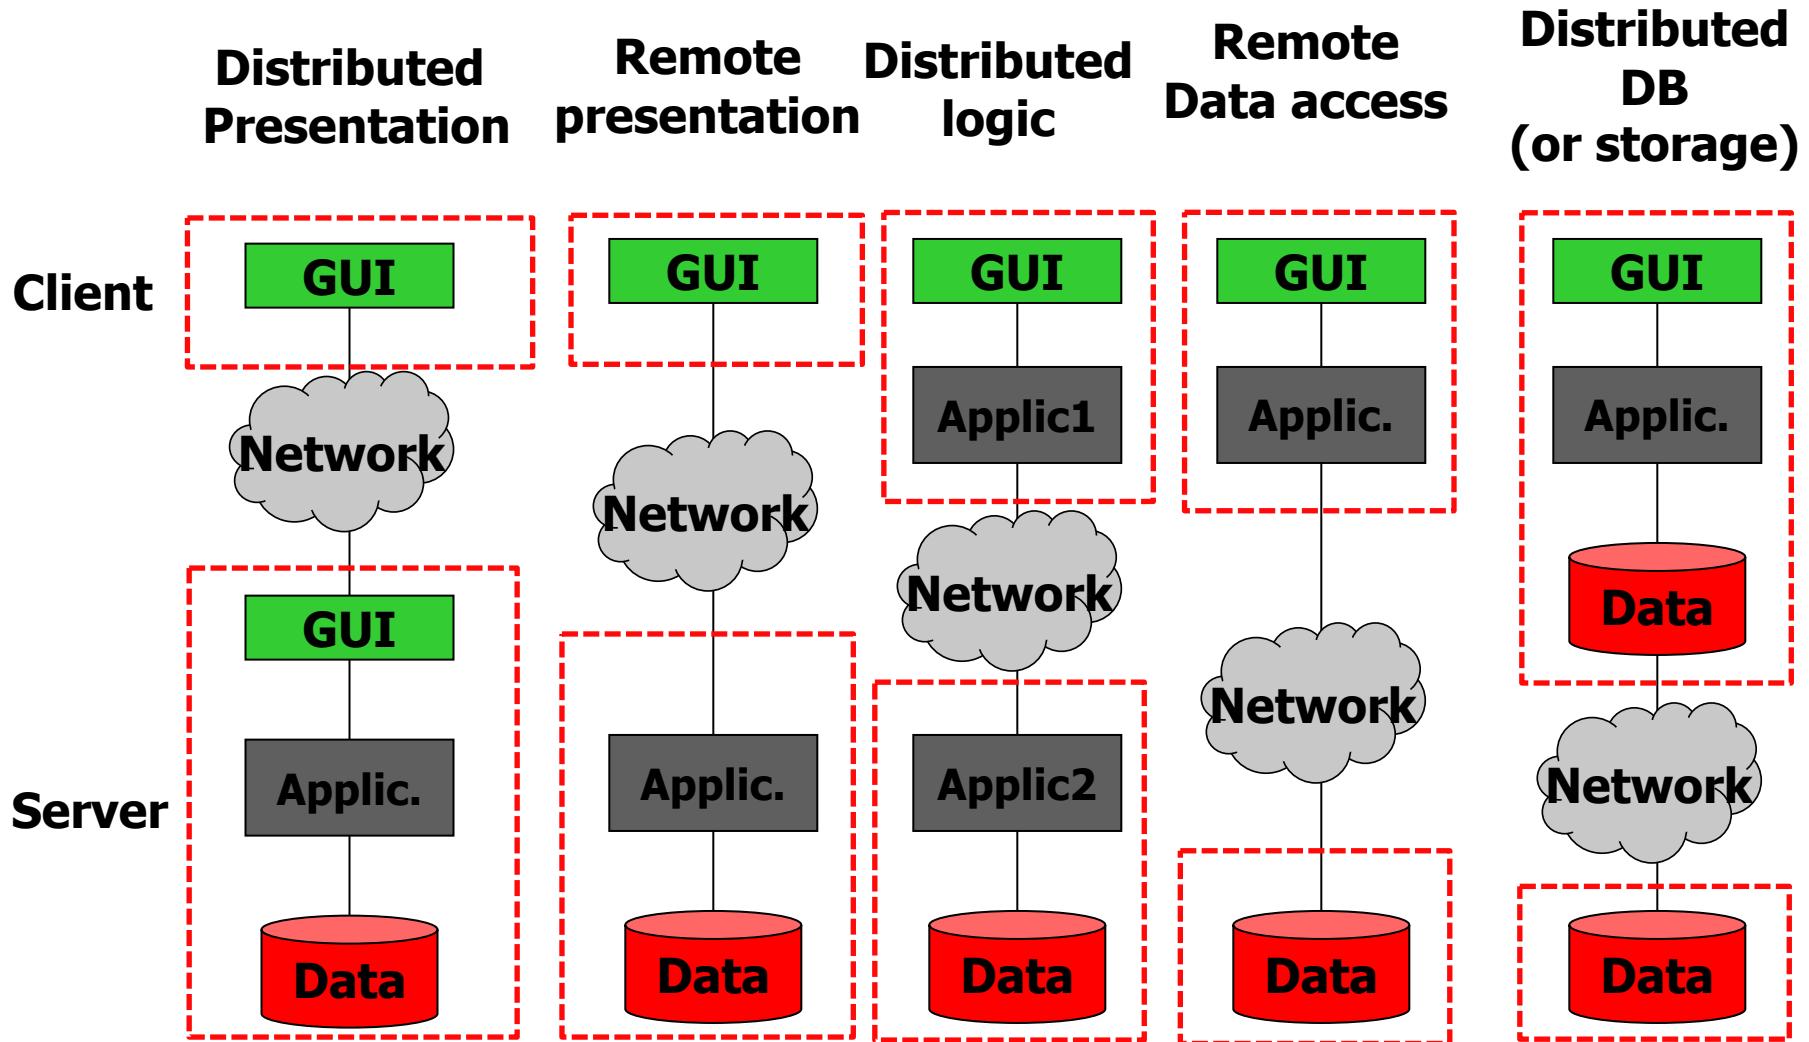


When to adopt a client-server architecture



- Remote communication
- Separation of concerns between the clients and the server
 - ▶ Note that a client-server architecture is a kind of layered architecture with two layers, also called tiers in this case
- Possibility to execute on heterogeneous computational units (virtual/physical)
 - ▶ E.g., servers on powerful machines (complex computations), clients on commodity machines

Client/Server = two-tier architecture



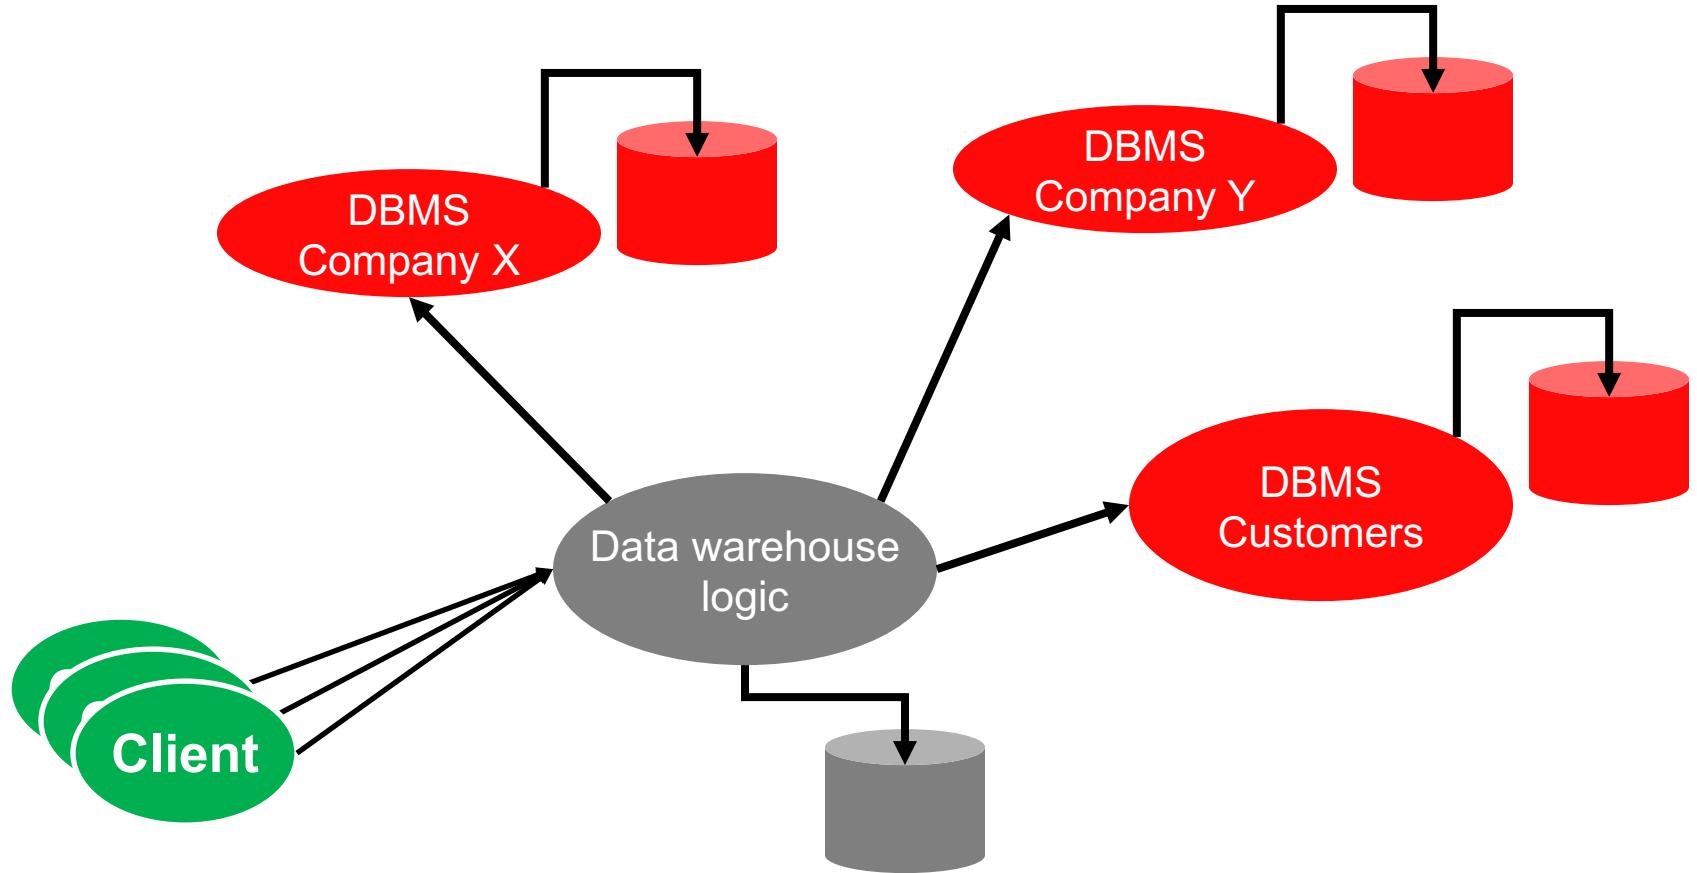


Three-tier architecture

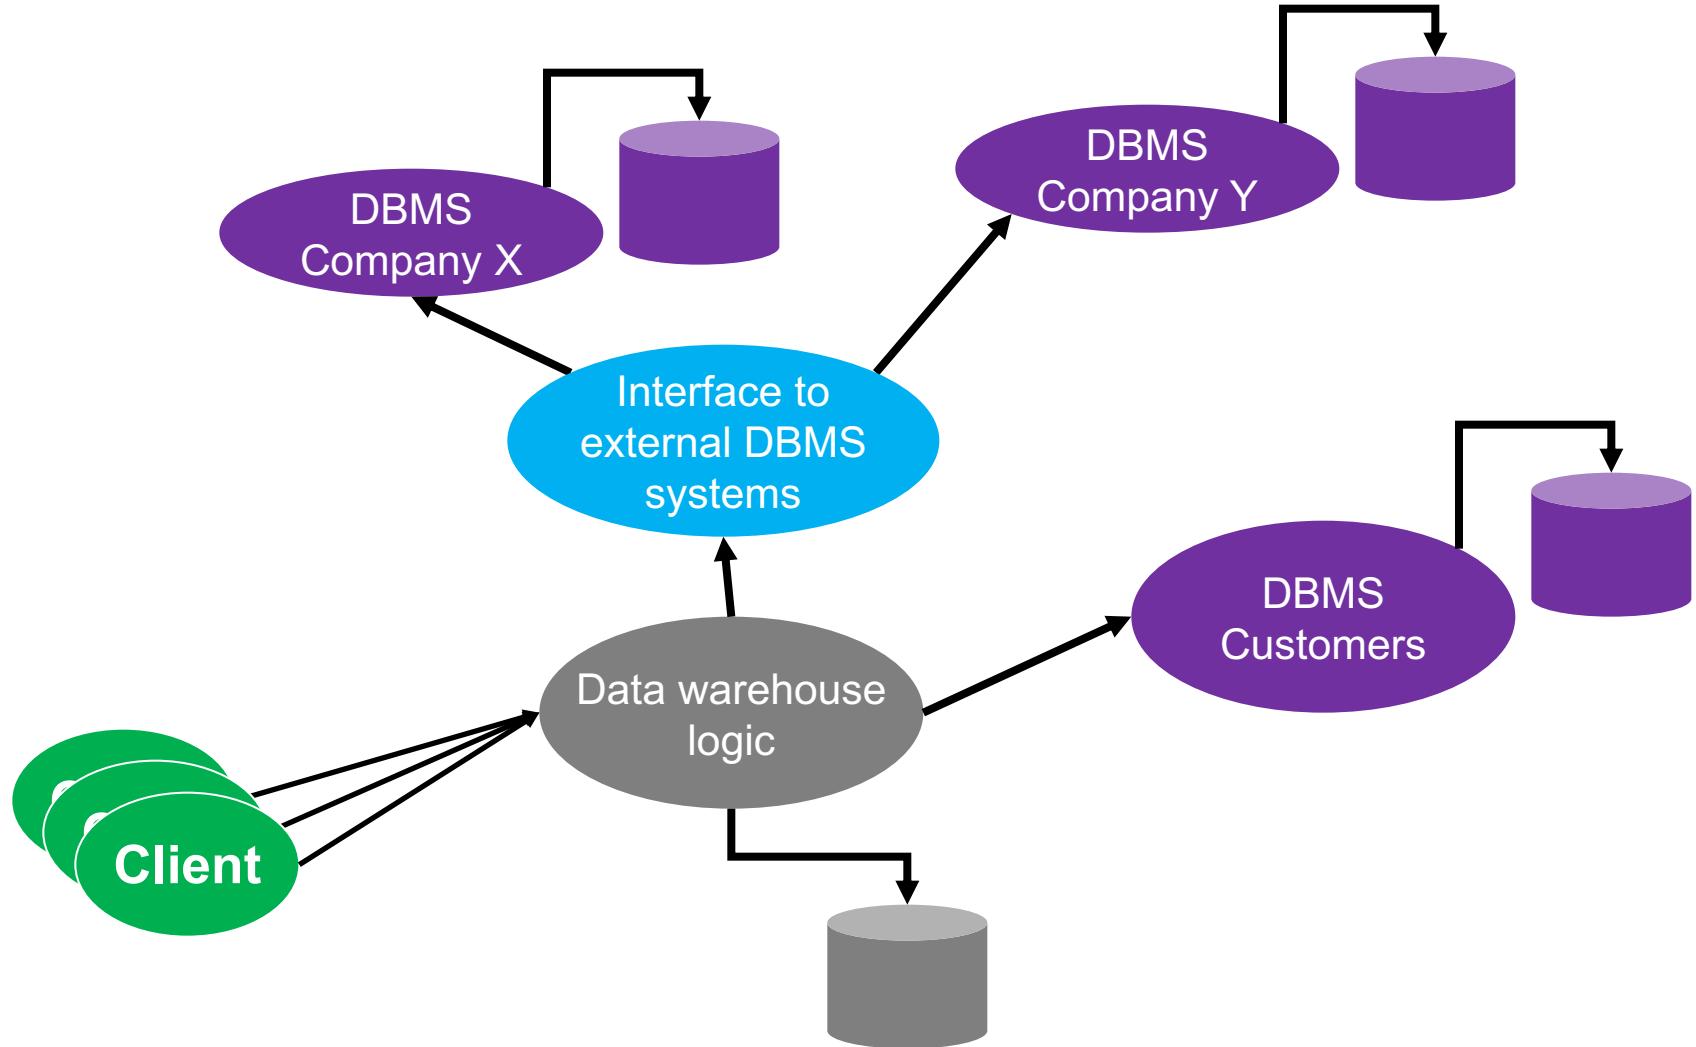
- In the typical scheme, the mid level has all the application logic
- Advantage
 - ▶ decoupling of logic and data, logic and presentation

An example

Data warehouse for a supermarket



From 3 to N levels





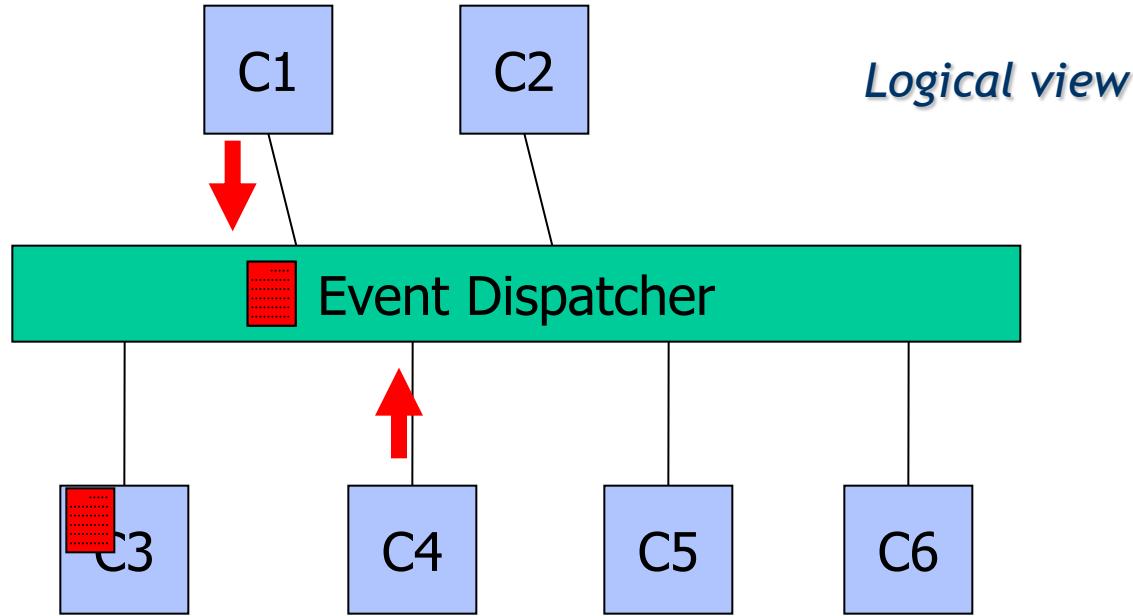
Relevant technologies

- JEE / Jakarta EE framework
 - ▶ <https://jakarta.ee/>
- Spring framework
 - ▶ <https://spring.io/>

Event-based systems

↑ Registration

■ Event



- Components can register to / send events
- Events are broadcast to all registered components
- No explicit naming of target component



Event-based paradigm

- Often called *publish-subscribe*
- Publish
 - ▶ event generation
- Subscribe
 - ▶ declaration of interest
- Different kinds of event languages possible



Event-based systems Pros & Cons

- + Very common in modern development practices
 - + E.g., continuous integration / deployment (such as GitHub Actions)
- + Easy addition/deletion of components
 - + Publishers/subscribers are decoupled
 - + The event dispatcher handles this dynamic set
- Potential scalability problems
 - The event dispatcher may become a bottleneck (under high workloads)
- Ordering of events
 - Not guaranteed, not straightforward



Other characteristics

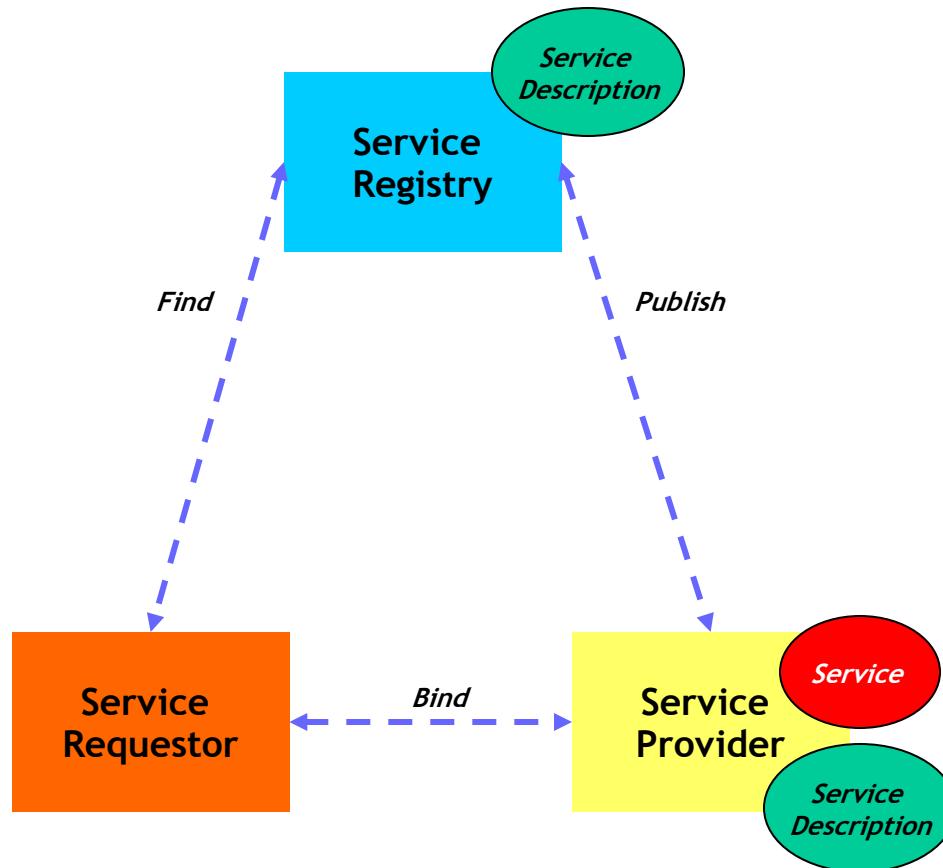
- Asynchrony
 - ▶ send and forget
- Reactive
 - ▶ computation driven by receipt of message
- Location/identity abstraction
 - ▶ destination determined by receiver, not sender
- Loose coupling
 - ▶ senders/receivers added without reconfiguration
- Flexible
 - ▶ one-to-many, many-to-one, many-to-many



Technologies

- Apache Kafka
 - ▶ <https://kafka.apache.org/>
- RabbitMQ
 - ▶ <https://www.rabbitmq.com/>

Service-oriented architecture (SOA)



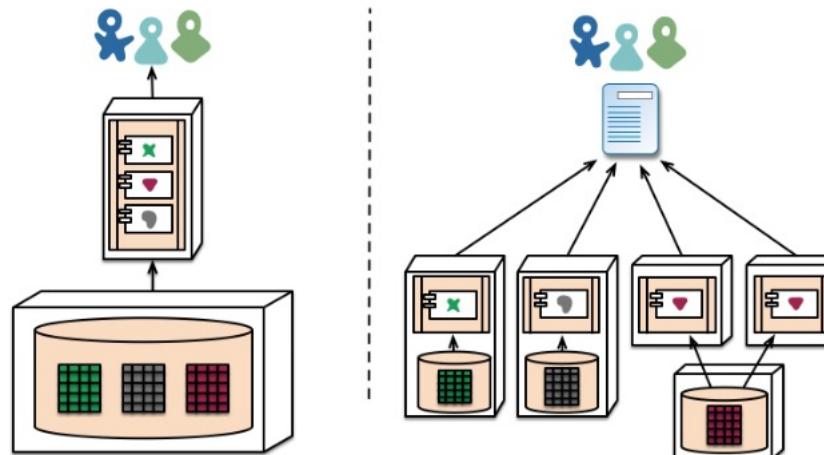


SOA Pros & Cons

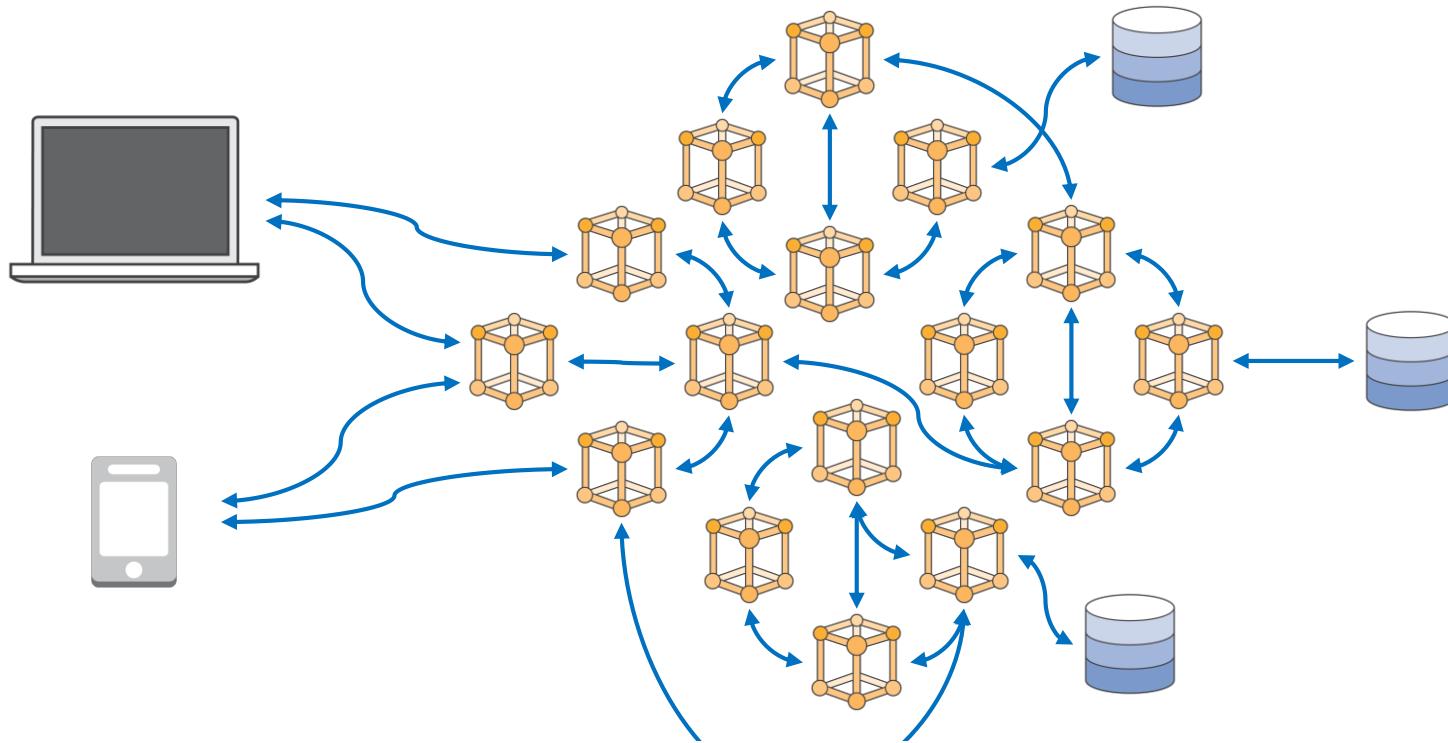
- Pros
 - ▶ Enables reuse of registries and providers across organizational boundaries
 - ▶ Relies on runtime discovery and dynamic binding
 - ▶ Services offered through a well-defined **Application Programming Interface (API)** → Interface documented in a machine-interpretable form
- Cons
 - ▶ Dynamic orchestration of discovered services is not trivial
 - A lot of glue code is needed
 - SOA systems are usually monolithic
 - ▶ SOAP (XML based protocol) considered too heavy

Microservice architectural style

- Evolution of SOA
 - ▶ <https://martinfowler.com/articles/microservices.html>
- Monolithic systems are decomposed into small specialized services and deal with a single “bounded context” in the target domain



Microservice architectural style



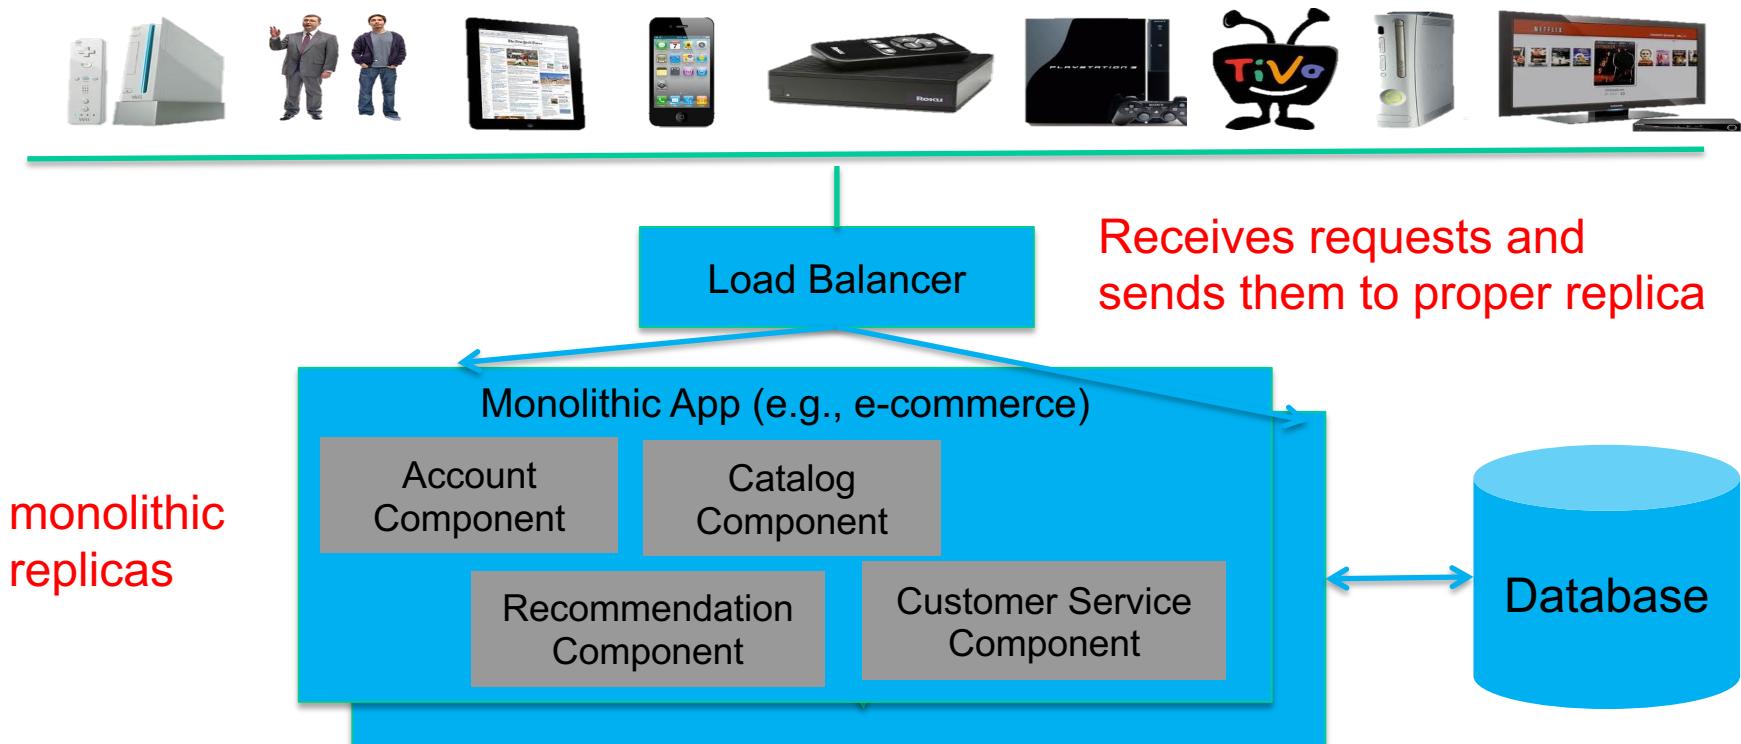
*“...the microservice architectural style is an approach to developing a single application as a suite of **small services**, each running in its own process and communicating with **lightweight mechanisms**, often an **HTTP resource API**”*

-- Martin Fowler

The starting point for microservices - monolithic architecture



Lots of users → high workload





Big problems (1/2)

- Issues with frequent deployments
 - ▶ Need to redeploy everything to change one component
 - ▶ Interrupt long running background jobs
 - ▶ Increase risk of failure
 - Overloaded containers
 - Obstacle to scaling development
 - ▶ Difficult fault localization
 - ▶ Difficult maintenance
 - Require long-term commitment to a technology stack
-

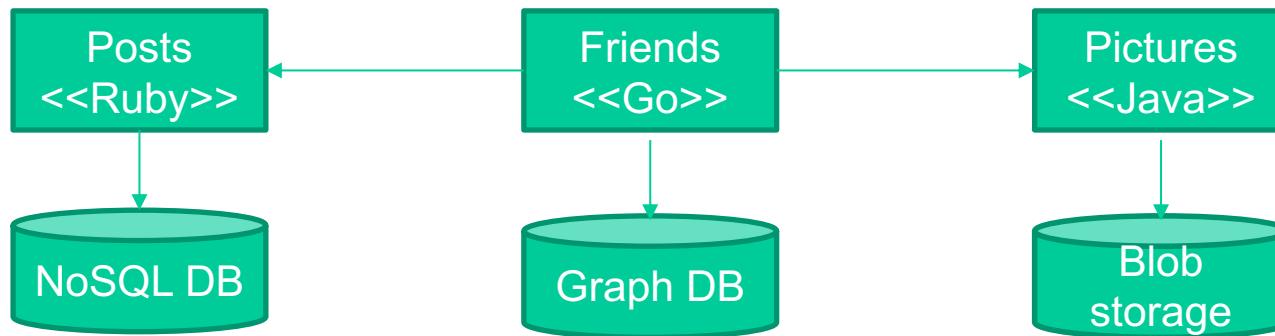


Big problems (2/2)

- Forbid fine-grained scaling strategies
 - ▶ It is not possible to scale “Product Catalog” up differently from “Customer Service”
- May lead to availability issues
 - ▶ [~2008] Netflix reported that a single minor mistake (i.e., missing “;”) brought down the whole platform for many hours

Microservices key benefits: technology heterogeneity

- Each service uses its own technology stack
 - ▶ Technology stack can be selected to fit the task best
 - E.g., Data analysis vs Video streaming
 - ▶ Teams can experiment with new technologies within a single microservice
 - E.g., we can deploy the two versions and do A/B testing
 - ▶ No unnecessary dependencies or libraries for each service





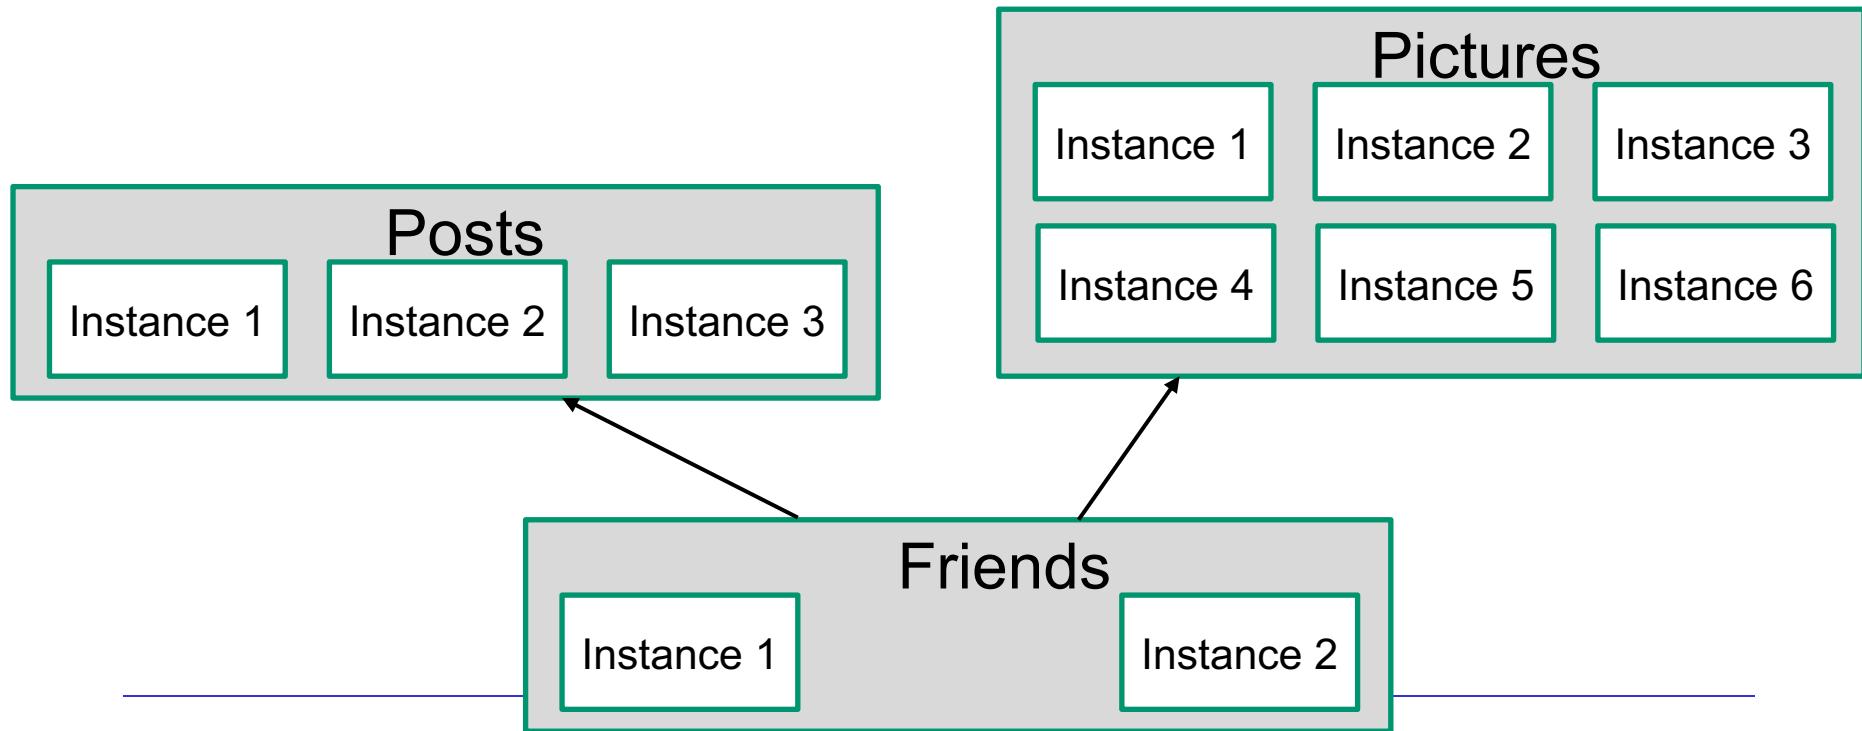
Microservices key benefits: resilience

- If a microservice fails the others can still work, possibly with a degraded functionality (until the failure is resolved)
- Note: microservices can introduce other types of failures related to network and communication (microservices systems are typically more “chatty”)



Microservices key benefits: scaling

- Each microservice can be scaled **independently**
- Identified bottlenecks can be addressed directly
 - ▶ Parts of the system that do not represent bottlenecks can remain simple and un-scaled



Microservices key benefits: independent codebases and deployment



- Each service has its own software repository
 - ▶ Favor distributed teams
 - ▶ No cross-dependencies between codebases
 - Tools work faster
 - ▶ Building, testing, refactoring, deployment takes seconds
 - Each service can be deployed independently
 - ▶ Independent CI/CD
 - Startup takes a smaller amount of time
-



Microservices key benefits: reuse and composability

- The functionality offered by a microservice can be used and reused in multiple contexts
 - ▶ E.g., authentication
- It is possible to compose multiple microservices in different ways



Microservices key benefit: replaceability

- Monolithic systems are often big and complex and their replacement is risky and costly
- Given its small size, replacing a microservice implementation is much easier
 - ▶ The team can develop a new implementation
 - ▶ As soon as the service is ready for operation, it can be moved in the operational environment
 - ▶ The previous service can be shutdown



Evolution at NETFLIX

2008

- Single DataCenter
- Everything in one WebApp

~2010

- AWS Cloud
- 100s of fine-grained Services

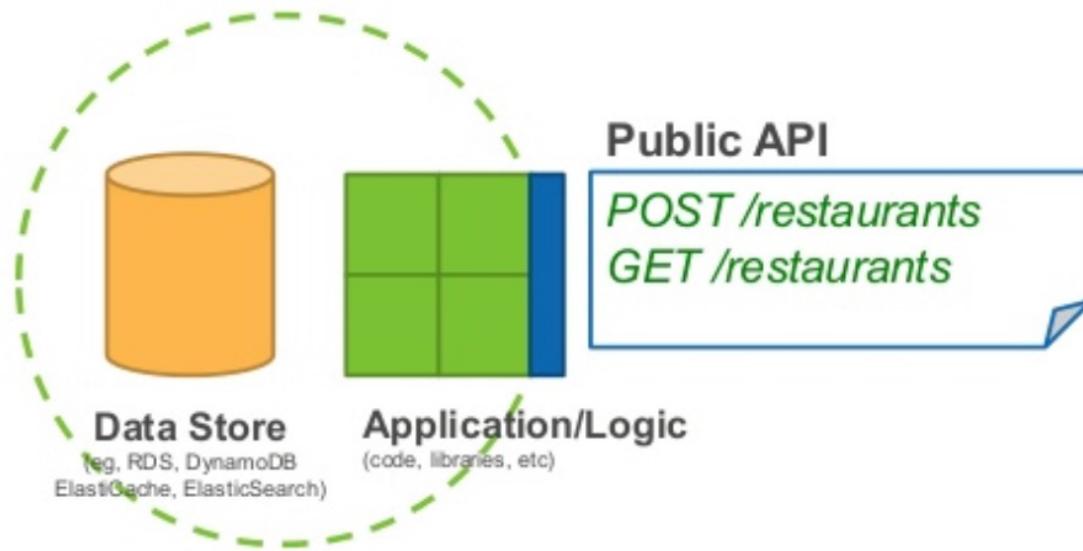


Current scale at NETFLIX

- 30% of the Internet traffic
 - 500+ microservices
 - 2+ billion API gateway requests daily
 - Each Gateway API call requires average six calls to backend services
 - Over 800 different client devices
-

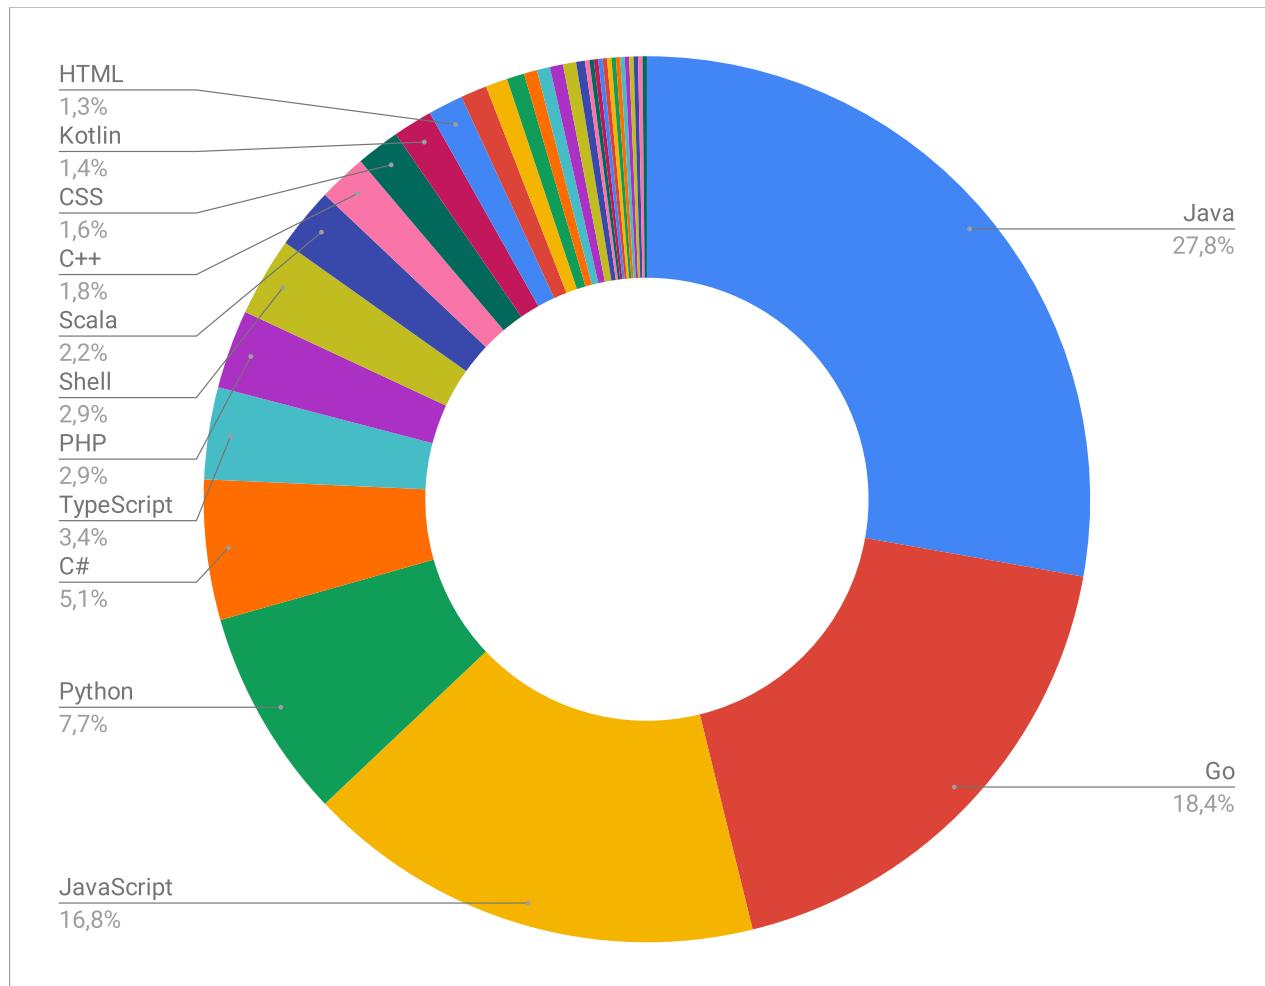


Anatomy of a microservice





Used languages

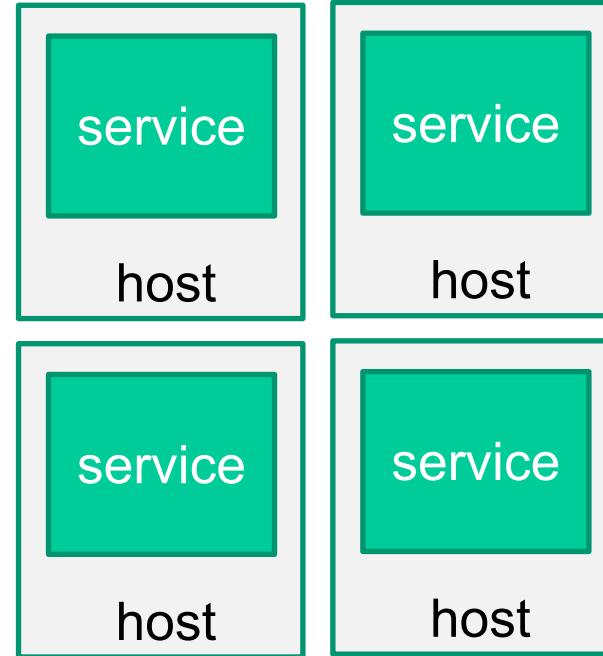
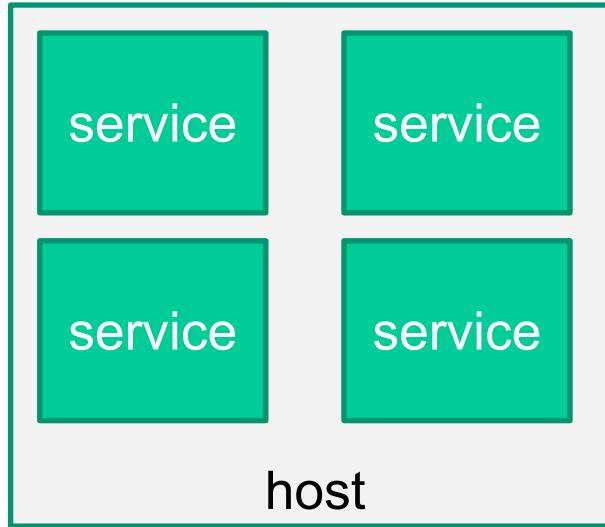




Critical issues with microservice architectures

- We deal with distributed systems
- Decomposition of monolithic systems require
 - ▶ Careful identification of microservices
 - ▶ Communication, integration and distributed transactions
 - ▶ Careful definition of the guiding principles and practices for the microservice architecture (microservices patterns/antipatterns)
 - ▶ Deployment
 - ▶ Testing
 - ▶ Monitoring

Multiservice vs single service per host deployment

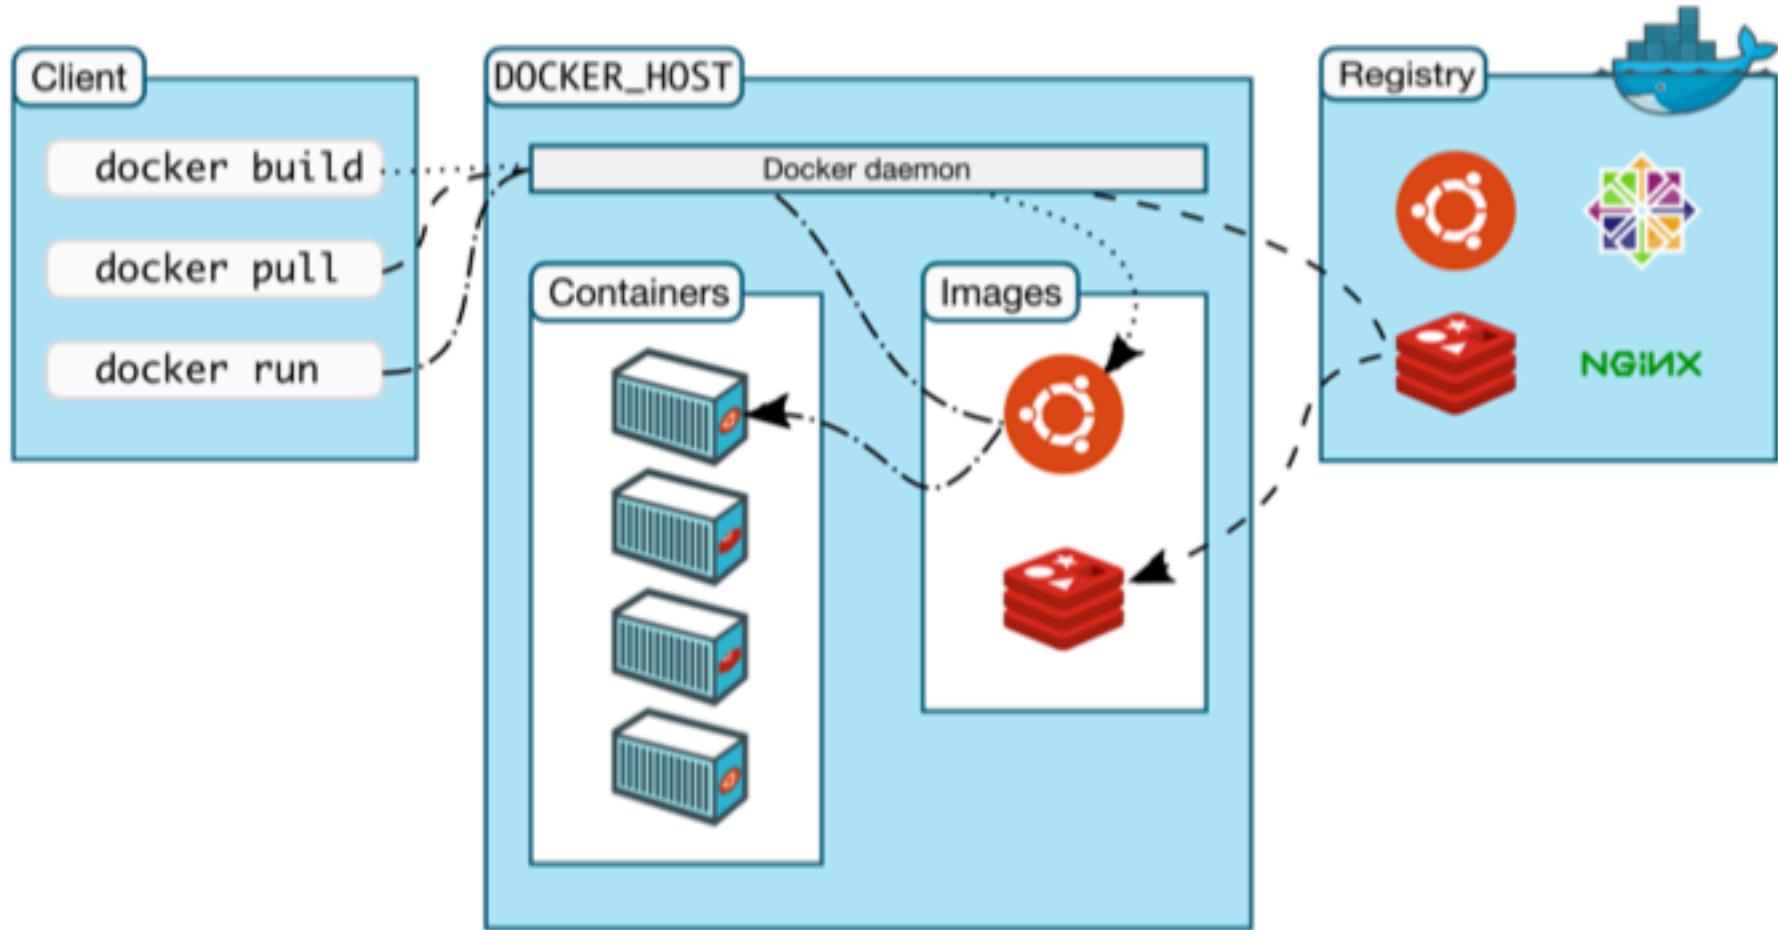


Requires more resources
Makes scalability easier and even

- Various mechanisms available for host virtualization



Possible container technology - Docker





... More Architectural Styles: Cloud Patterns

<http://www.cloudcomputingpatterns.org>



Classification of patterns

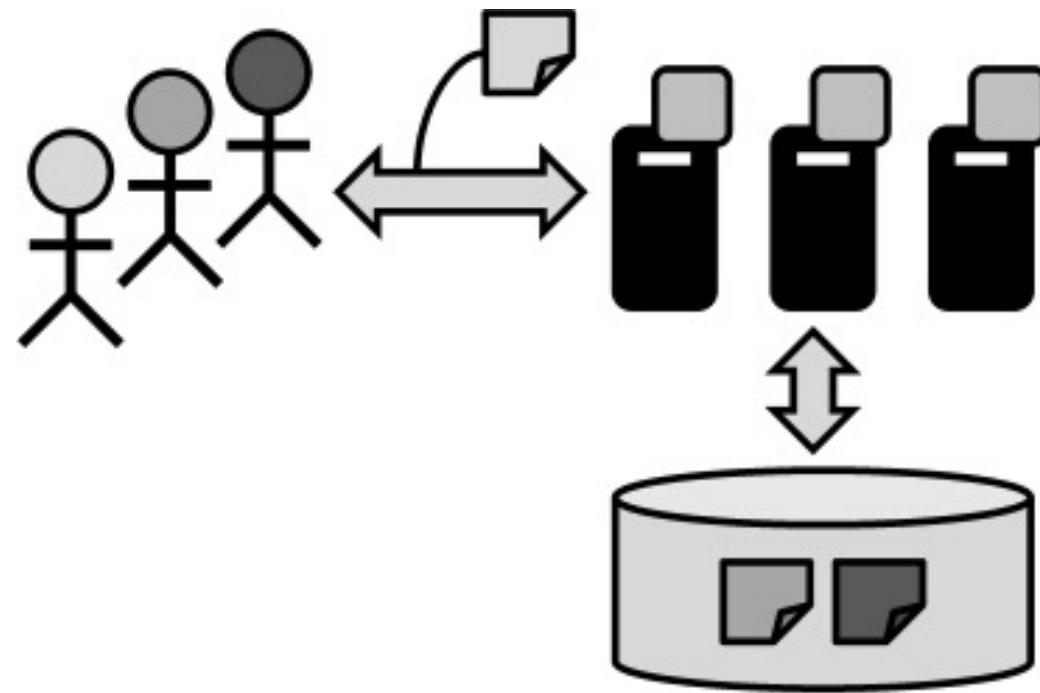
- Basic Architectural Patterns
 - ▶ Composite Application
 - ▶ Loose Coupling
 - ▶ Stateless Component
 - ▶ Idempotent Component
- Elasticity Patterns
 - ▶ Map Reduce
 - ▶ Elastic Component
 - ▶ Elastic Load Balancer
 - ▶ Elastic Queue
- Availability Patterns
 - ▶ Watchdog
 - ▶ Update Transition
- Multi-Tenancy Patterns
 - ▶ Single Instance Component
 - ▶ Single Configurable Instance Component
 - ▶ Multiple Instance Component



Stateless components

- Context
 - ▶ A componentized application subsumes components that can fail.
 - Challenges
 - ▶ Cloud resources can display low availability.
 - ▶ Component instances can be added and removed regularly when the demand changes.
 - Solution
 - ▶ Implement the components in a way that they do not contain any internal state, but completely rely on external persistent storage.
-

Stateless components



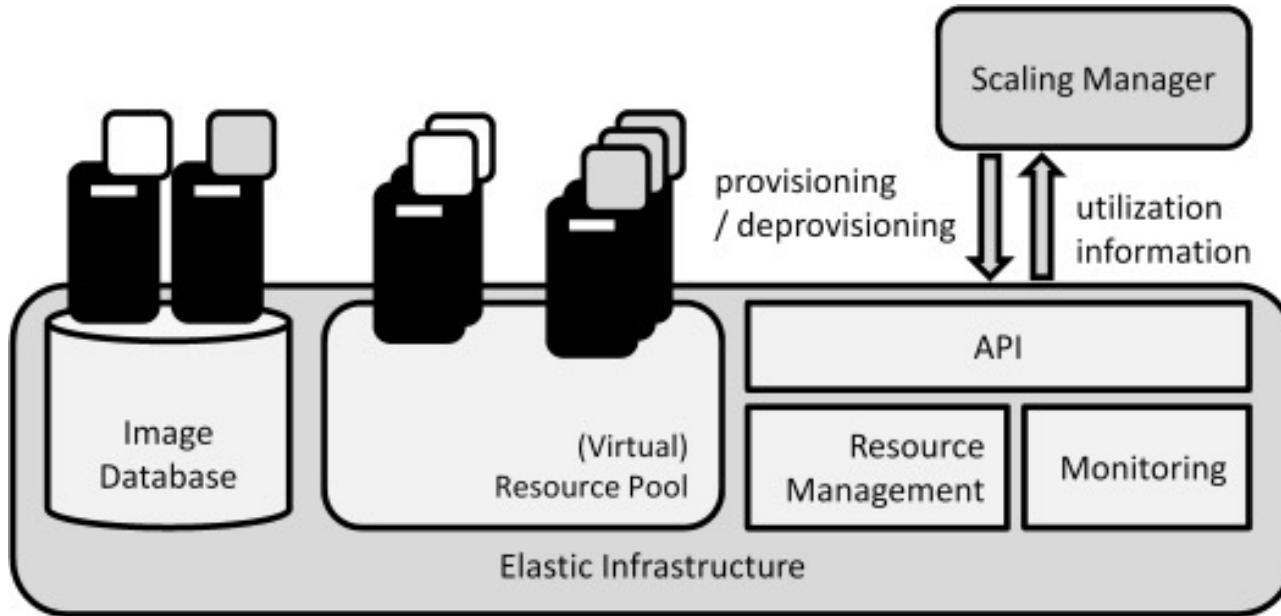
- No data is lost if an instance fails
- Multiple components have a common internal state
- The scalability of the central data store becomes the major challenge
 - ▶ Distributed data stores



Elastic component

- Context
 - ▶ A componentized application uses multiple compute nodes provided by an *elastic infrastructure*.
 - Challenges
 - ▶ To benefit from the elastic infrastructure, the management process to scale-out a componentized application has to be automated.
 - ▶ Manual resource scaling would not support pay-per-use nor quickly changing workload
 - Solution
 - ▶ Monitor the utilization of compute nodes that host application components and automatically adjust their numbers using the provisioning functionality provided by the elastic infrastructure.
-

Elastic component



- If the utilization of compute nodes exceeds a specified threshold, additional hosting components are provisioned that contain the same application component.
- If the components are implemented as *stateless components*, the operations for adding and removal of components are significantly simplified.

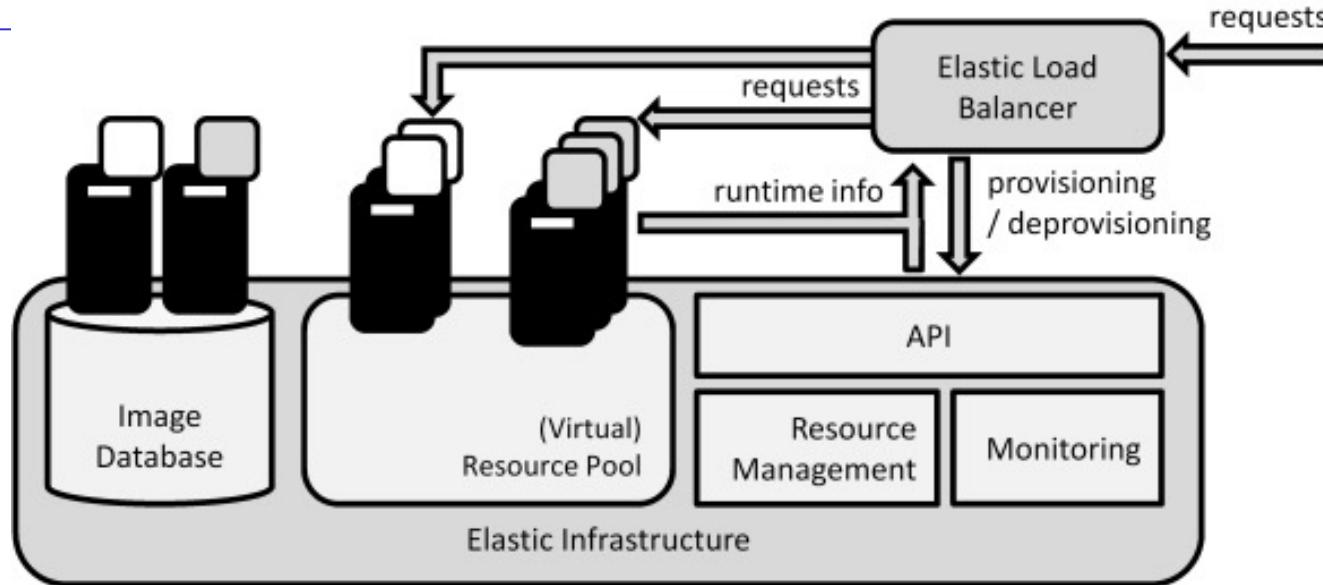


Elastic load balancer

- Context
 - ▶ A componentized application uses multiple compute nodes provided by an elastic infrastructure.
 - Challenges
 - ▶ As for elastic component
 - ▶ Service requests are a good indicator of workload and therefore shall be used as a basis for scaling decisions
 - Solution
 - ▶ Use an elastic load balancer that determines the amount of required resources based on numbers of requests and provisions the needed resources accordingly using the elastic infrastructure's API
-



Elastic load balancer



- The number of requests that can be handled by a computing node is a crucial design parameter
- It has to be adjusted at run time
- Information about the time needed to provision new compute nodes can also be necessary to deduct effective scaling actions



Assignments for the next class

- Watch the videos you find here
 - ▶ On software design descriptions and principles:
https://polimi365-my.sharepoint.com/:v/g/personal/10143828_polimi_it/EaE5ix1izNxGiLLtzBsr7GABI6M_BqvtBPeH4qtuVMOnhQ?e=hydFes
 - ▶ On the design process: https://polimi365-my.sharepoint.com/:v/g/personal/10143828_polimi_it/EcbUpgqCblMiw2hIZIUFSAff1zSZCjXuYrgt1cQfntw?e=ebYZC9
- They will be used as a basis for discussion in class next time